

SOFT3410: Concurrency for Software Developers

Transactional memory

Lecturer: Martin McGrane
School of Computer Science



THE UNIVERSITY OF
SYDNEY



Concurrent Programming is Hard

- Hard to make **safe** and **live**
 - Safe: *where nothing bad (e.g., crash) happens*
 - Live: *where something good (that we want) eventually happens*
- The human mind is usually trained to be sequential
 - Concurrent specifications
 - Non-deterministic executions

Safety

```
public class BankAccount {  
    int balance = 0;  
  
    void deposit(int amount) {  
  
        balance += amount;  
  
    }  
}
```

Safety

```
public class BankAccount {  
    int balance = 0;  
  
    void deposit(int amount) {  
  
        balance += amount;  
  
    }  
}
```

No concurrency control: race!

Liveness

› Problem

```
public class BankAccount {  
    int balance = 0;  
  
    void deposit(int amount) {  
        balance += amount;  
    }  
}
```

```
temp = balance;  
balance = temp + amount;
```

Inconsistency

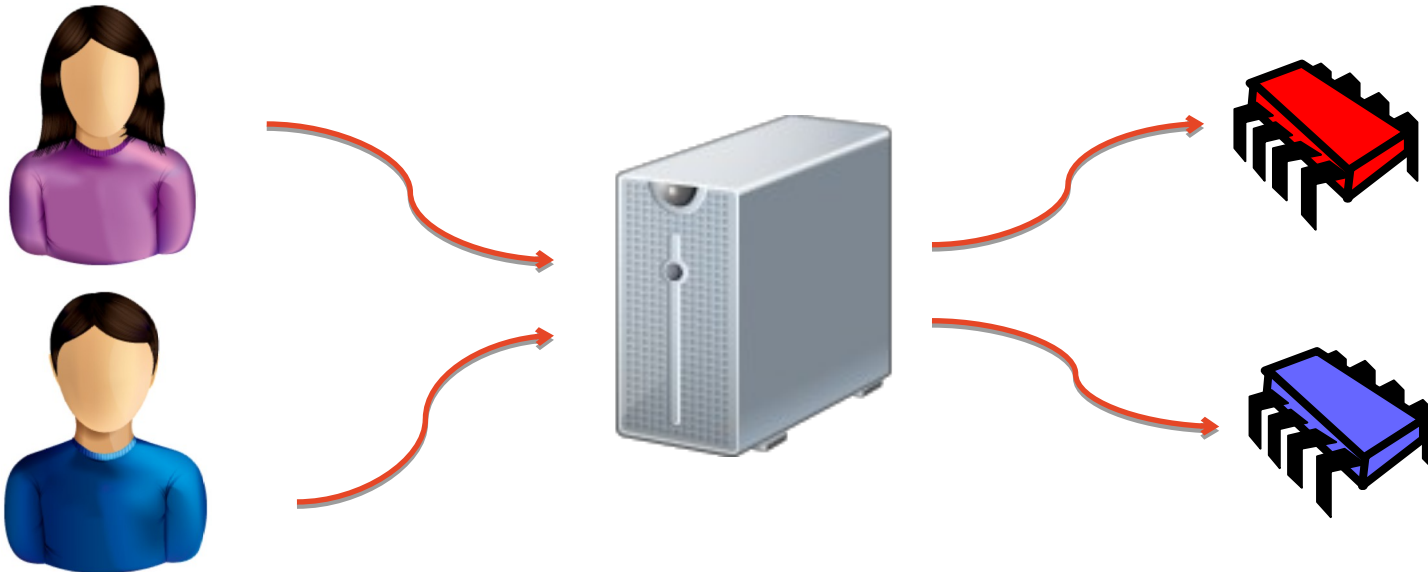
- › Let Alice and Bob deposit on the same account **a** concurrently:

```
a.deposit(20AUD) ;
```

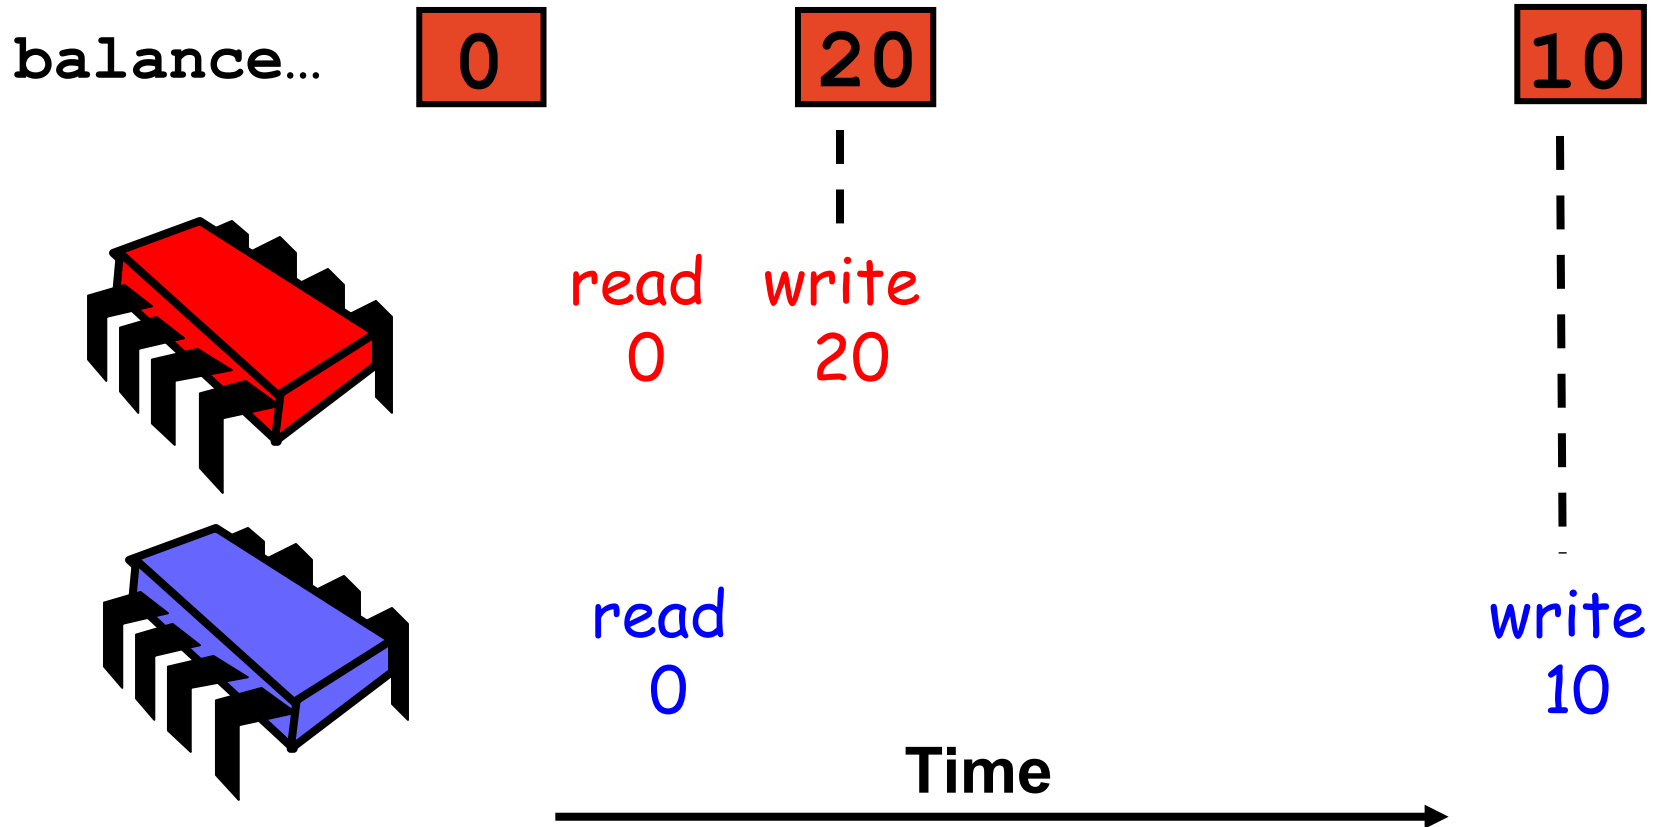
//

```
a.deposit(10AUD) ;
```

- › The multi-threaded server handles the two requests concurrently:



Not so good...



Safety

› Challenge

```
public class BankAccount {  
    int balance = 0;  
  
    void deposit(int amount) {  
  
        temp = balance;  
        balance = temp + amount;  
  
    }  
}
```

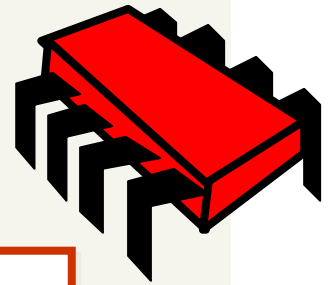
Make these steps indivisible, linearisable

Safety

› Hardware solution

```
public class BankAccount {  
    int balance = 0;  
  
    void deposit(int amount) {  
  
        balance += amount;  
  
    }  
}
```

Read Modify Write Instruction



Safety

› Software solution?

```
public class BankAccount {  
    int balance = 0;  
  
    void deposit(int amount) {  
        synchronized(this) {  
            temp = balance;  
            balance = temp + amount;  
        }  
    }  
}
```

Safety

› Software solution?

```
public class BankAccount {  
    int balance = 0;  
  
    void deposit(int amount) {  
        synchronized(this) {  
            temp = balance;  
            balance = temp + amount;  
        }  
    }  
}
```

Synchronised block

Safety

› Software solution?

```
public class BankAccount {  
    int balance = 0;  
  
    void deposit(int amount) {  
        synchronized(this) {  
            temp = balance;  
            balance = temp + amount;  
        }  
    }  
}
```

Mutual Exclusion

Safety

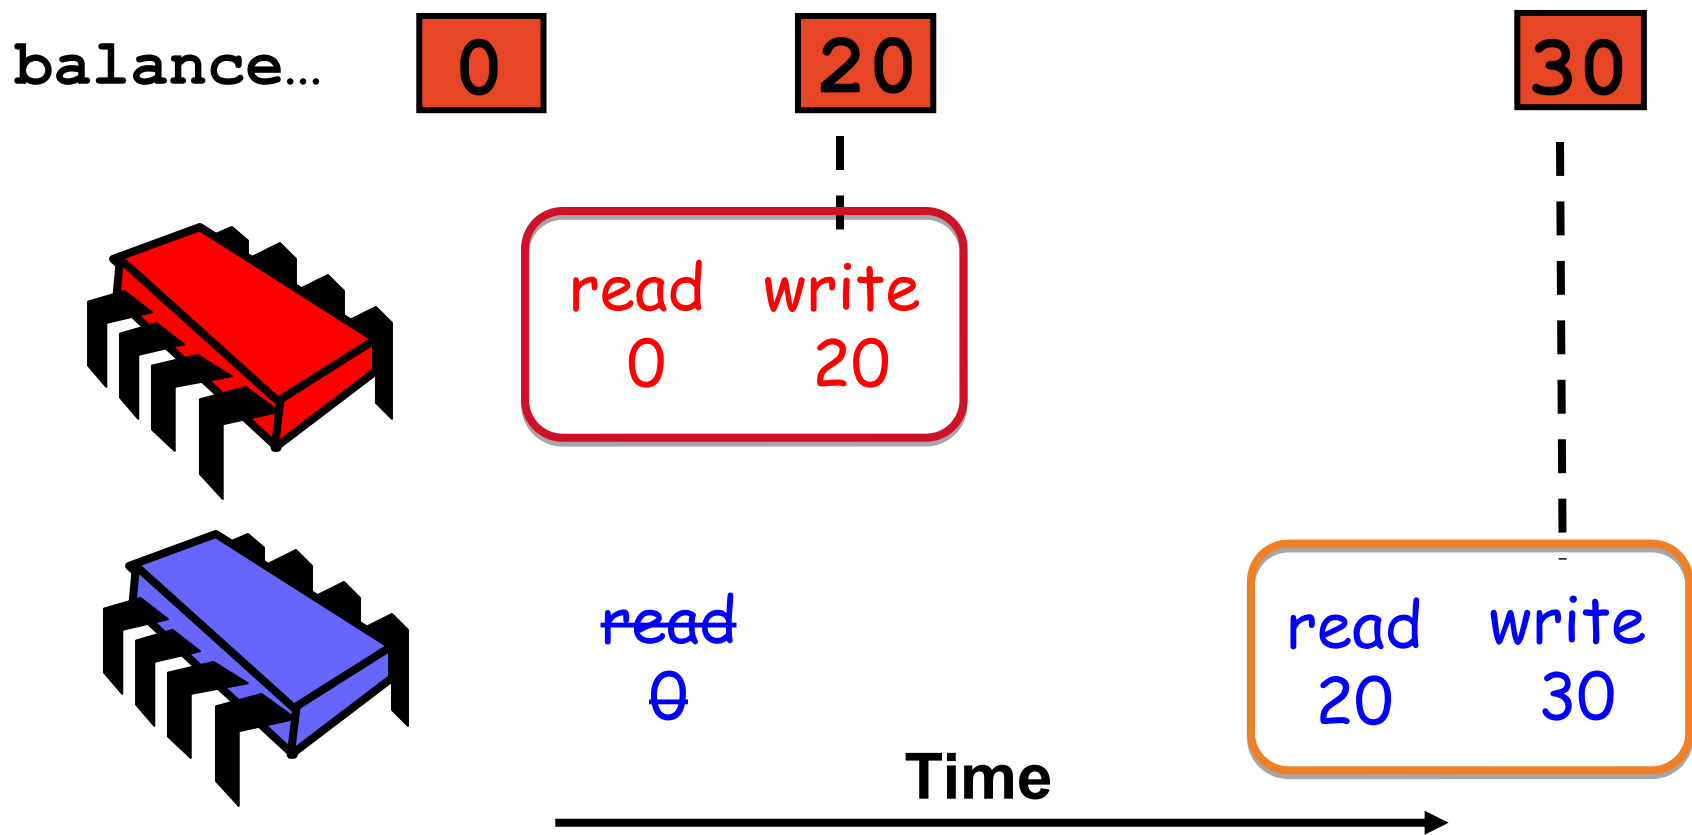
- › Multiple threads **cannot enter a synchronised block** with the same argument at the **same time**. The argument represents the resource that can be used by a single thread at a time.

```
a.deposit(x) ;
```

```
//
```

```
a.deposit(y) ;
```

Not so bad...



Concurrent Programming is Hard

- Hard to make **safe** and **live**
 - Safe: *where nothing bad (e.g., crash) happens*
 - Live: *where something good (that we want) eventually happens*
- The human mind is usually trained to be sequential
 - Concurrent specifications
 - Non-deterministic executions

We obtained a safe program

Concurrent Programming is Hard

- Hard to make **safe** and **live**
 - Safe: *where nothing bad (e.g., crash) happens*
 - Live: *where something good (that we want) eventually happens*
- The human mind tends to be sequential
 - Concurrent specifications
 - Non-deterministic executions

But is the program also live?

Liveness

› A slightly more complex problem...

```
void deposit(...) { synchronized(this) { ... } }  
void withdraw(...) { synchronized(this) { ... } }  
int balance(...) { synchronized(this) { ... } }
```

```
void transfer(account from, int amount) {
```

No concurrency control: race!

```
    if (from.balance() >= amount) {  
        from.withdraw(amount);  
        this.deposit(amount);  
    }
```

```
}
```

Liveness

```
void deposit(...) { synchronized(this) { ... } }
void withdraw(...) { synchronized(this) { ... } }
int balance(...) { synchronized(this) { ... } }

void transfer(account from, int amount) {

    synchronized(this) {
        if (from.balance() >= amount) {
            from.withdraw(amount);
            this.deposit(amount);
        }
    }

}
```

Liveness

```
void deposit(...) { synchronized(this) { ... } }
void withdraw(...) { synchronized(this) { ... } }
int balance(...) { synchronized(this) { ... } }

void transfer(account from, int amount) {

    synchronized(this) {
        if (from.balance() >= amount) {
            from.withdraw(amount);
            this.deposit(amount);
        }
    }

}
```

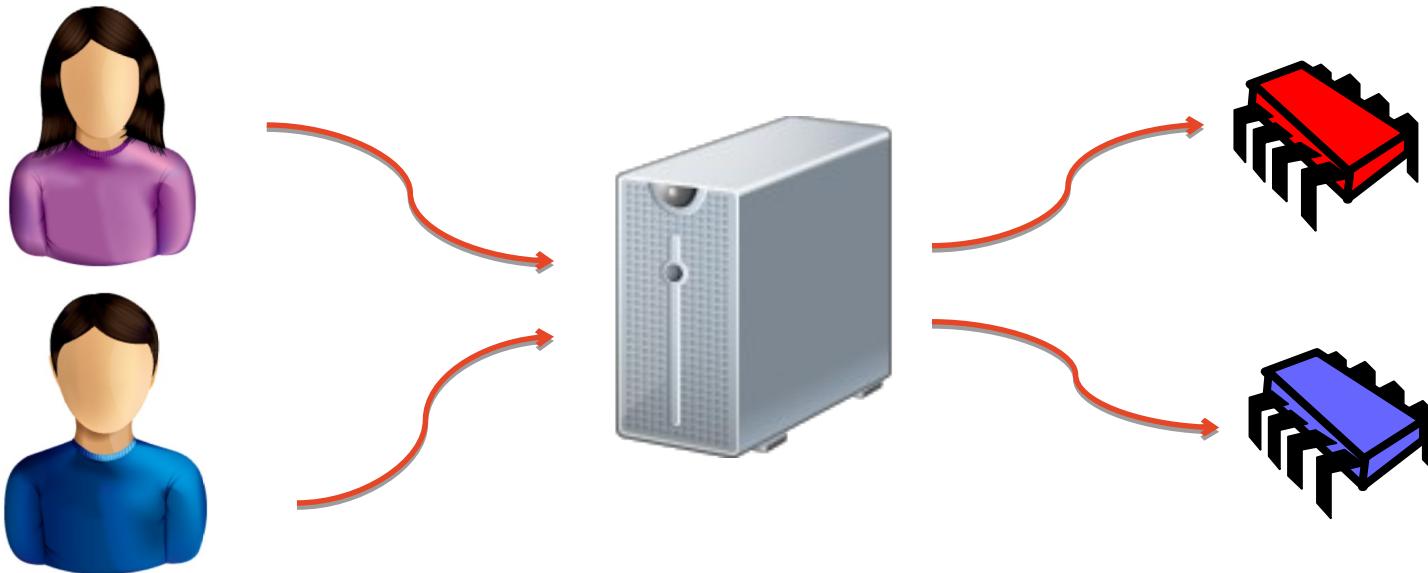
Race!

Inconsistency

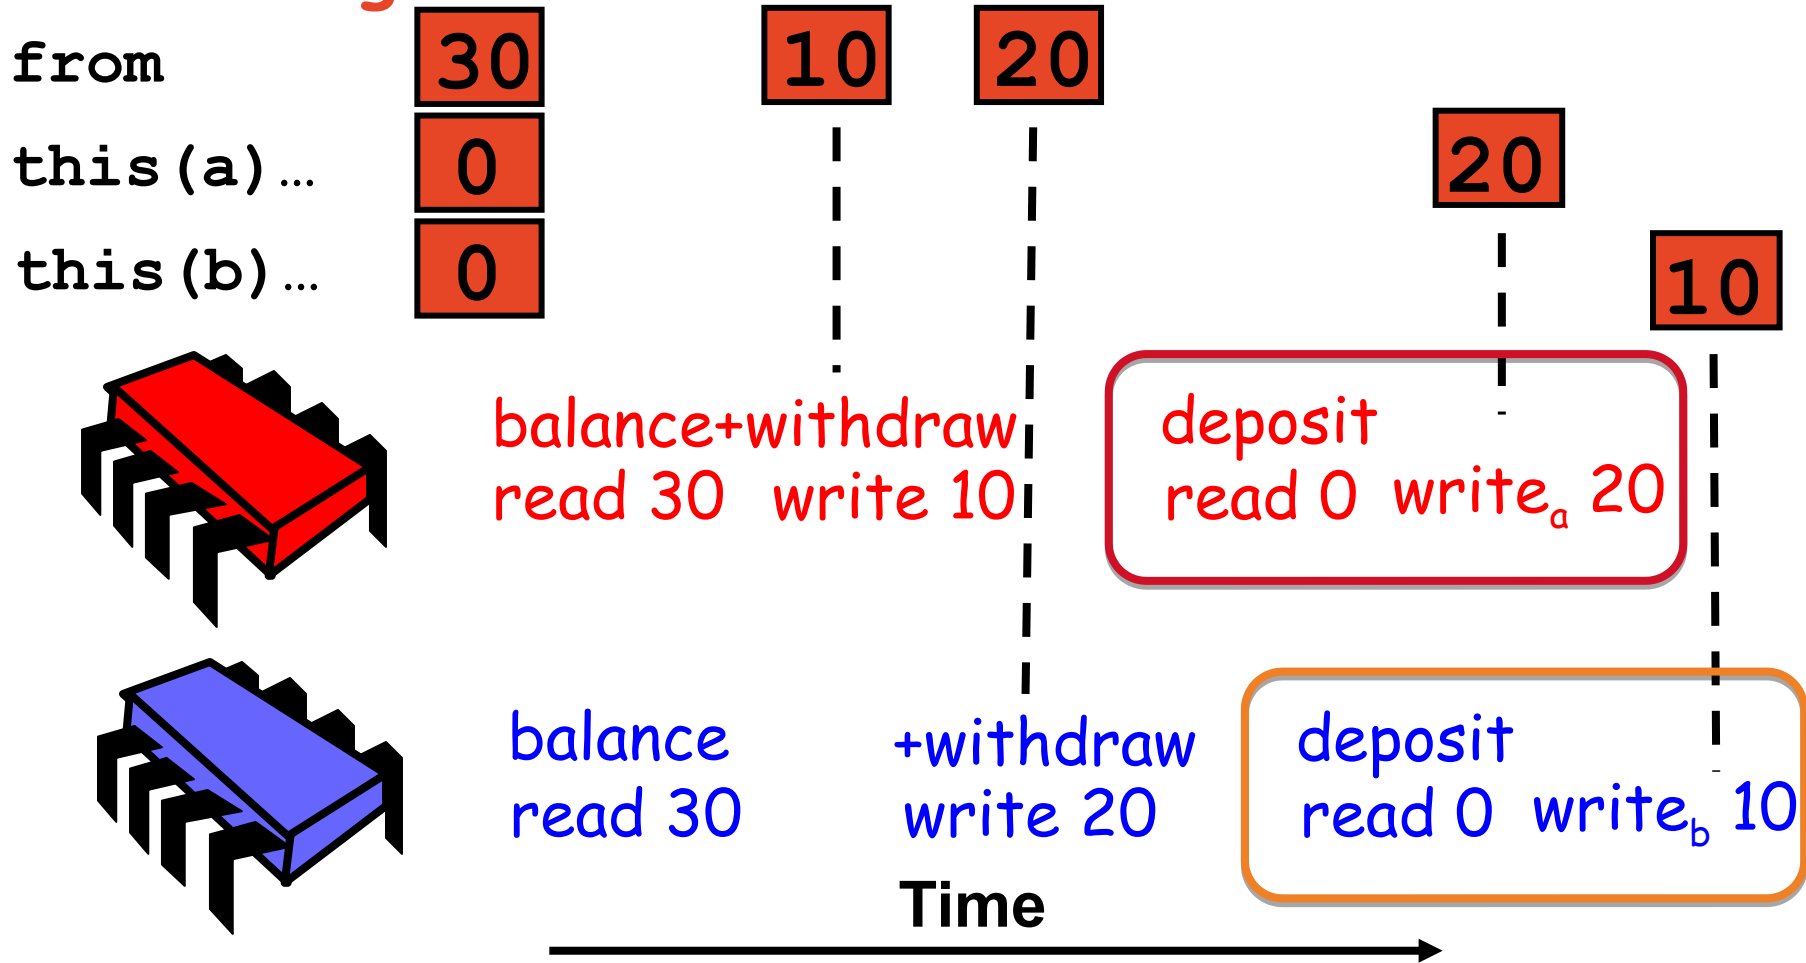
- › Let Alice and Bob between the same accounts concurrently:

```
a.transfer(b, 20AUD) ; // a.transfer(b, 10AUD) ;
```

- › The multi-threaded server handles the two requests concurrently:



Not so good...



Liveness

```
void deposit(...) { synchronized(this) { ... } }
void withdraw(...) { synchronized(this) { ... } }
int balance(...) { synchronized(this) { ... } }

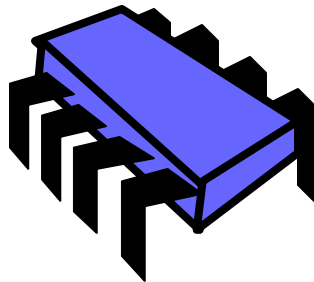
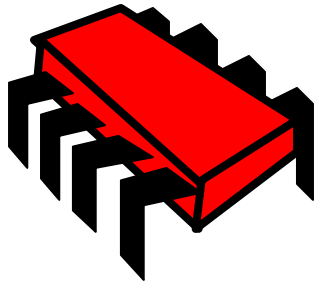
void transfer(account from, int amount) {
    synchronized(this) {
        synchronized(from) {
            if (from.balance() >= amount) {
                from.withdraw(amount);
                this.deposit(amount);
            }
        }
    }
}
```

Better...

from

this (a)...

this (b)...



30
0
0

10

20

10

10

withdraw deposit
r30 w10 r0 w20

withdraw deposit
r10 w0 r0 w10

Time



Liveness

```
void deposit(...) { synchronized(this) { ... } }
void withdraw(...) { synchronized(this) { ... } }
int balance(...) { synchronized(this) { ... } }

void transfer(account from, int amount) {
    synchronized(this) {
        synchronized(from) {
            if (from.balance() >= amount) {
                from.withdraw(amount);
                this.deposit(amount);
            }
        }
    }
}
```

Deadlock!

Deadlock

- › Let two threads transfer between accounts **a** and **b** in opposite order:

```
a.transfer(b, x) ;
```

```
//
```

```
b.transfer(a, y) ;
```

Deadlock

- › Let two threads transfer between accounts **a** and **b** in opposite order:

```
a.transfer(b, x) ;    //    b.transfer(a, y) ;
```

- › Each thread may successfully acquire the first lock

```
...  
synchronized(a) {  
    synchronized(b) {  
        ...  
    }  
}
```

```
...  
synchronized(b) {  
    synchronized(a) {  
        ...  
    }  
}
```

- › If so, they will never be able to acquire the second lock!

Deadlock



Deadlock



Deadlock

Consequences of deadlocks: examples in operating systems

› Google FS:

- To make a file rename atomic when moving a file from directory **a** to directory **b**, locks must be acquired on **a** and **b**.
- Concurrent renames moving from **a** to **b** and from **b** to **a** may deadlock.
- Avoided by requiring lexicographical ordering for lock acquisition



› Linux kernel:

- `linux/mm/filemap.c` starts with 50 lines of commented code to explain the order in which locks are acquired.
- A kernel programmer must first understand this before locking anything, otherwise risking to deadlock.



Deadlock

- Multiple locks acquired in arbitrary order may lead to deadlock
- Multiple locks acquired in a single shared order are safe from deadlock
 - Lexicographical order
 - Memory address order

Transactional Memory



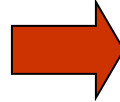
THE UNIVERSITY OF
SYDNEY

Transactional memory

An appealing alternative:

```
void deposit(int x) {  
    synchronized(this) {  
        int tmp = balance;  
        tmp += x;  
        balance = tmp;  
    }  
}
```

Lock acquire/release



```
void deposit(int x) {  
    atomic {  
        int tmp = balance;  
        tmp += x;  
        balance = tmp;  
    }  
}
```

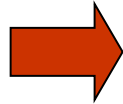
(As if) no interleaved computation

Easier-to-use primitive
(but harder to implement)

Transactional memory

Compiler-based instrumentation:

```
void deposit(int x) {  
    atomic {  
        int tmp = balance;  
        tmp += x;  
        balance = tmp;  
    }  
}
```



```
void deposit(int x) {  
    tx-start();  
    int tmp = tx-read(balance);  
    tmp += x;  
    tx-write(balance, tmp);  
    tx-commit();  
}
```

A safe and live (deadlock-free) program

Exploits the TM wrappers

A compiler supporting TM takes care of the instrumentation

Transactional memory

```
void deposit(...) { atomic { ... } }  
void withdraw(...) { atomic { ... } }  
int balance(...) { atomic { ... } }  
  
void transfer(account from, int amount) {  
  
    if (from.balance() >= amount) {  
        from.withdraw(amount);  
        this.deposit(amount);  
    }  
  
}
```

No concurrency control: race!

Transactional memory

```
void deposit(...) { atomic { ... } }  
void withdraw(...) { atomic { ... } }  
int balance(...) { atomic { ... } }  
  
void transfer(account from, int amount) {  
    atomic {  
        if (from.balance() >= amount) {  
            from.withdraw(amount);  
            this.deposit(amount);  
        }  
    }  
}
```

Correct and enables parallelism!

Transactional memory

Ensures safety (atomicity) and liveness (deadlock-freedom)

- Transactions: a simple paradigm

**A sequence of instructions,
executed atomically**

- Software transactions are good for:
 - Software engineering (simple programming, avoid races & deadlocks, composability)
 - Performance (when no conflict, high parallelism and no locking overhead)

A “universal” synchronization construct

Don't care how transactions are implemented!

Transactional memory

Implementation example:

- **tx-read(x) / tx-write(x, v) :**
 - if unlocked, then acquire lock of x
 - else **tx-abort()**
- **tx-commit() :**
 - release all locks
 - cleanup

Transactional memory

Implementation example (continued):

```
tx-write(x,v) {  
    if (!<x,v'> in rSet U wSet)  
        while(!CAS(lock(x),unlocked,locked))  
            tx-abort()  
    wLog = wLog U {<x,store(x,v)>}  
    wSet = wSet U {<x,v>}  
    return ok  
}
```

Transactional memory

Implementation example (continued):

```
tx-read(x) {  
    if (!<x,v> in rSet U wSet)  
        while(!CAS(lock(x),unlocked,locked))  
            tx-abort()  
    v = load(x)  
    rSet = rSet U {<x,v>}  
    return v  
}
```

Transactional memory

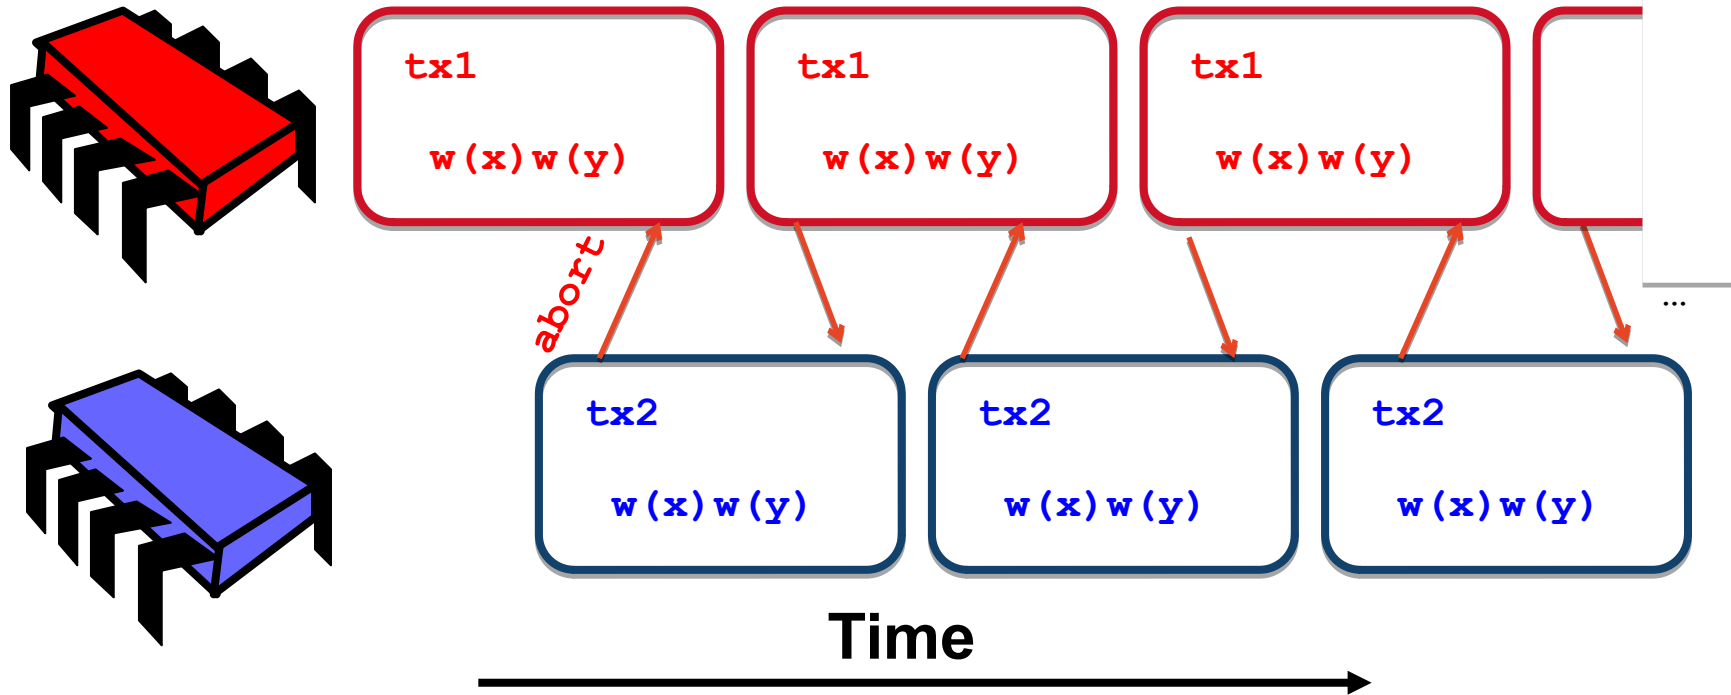
Implementation example (continued):

```
tx-commit() {  
    for (x in rSet U wSet) unlock(x)  
    empty rSet and wSet  
}
```

```
tx-abort() {  
    rollback(wLog)  
    for (x in rSet U wSet) unlock(x)  
    empty rSet and wSet  
}
```

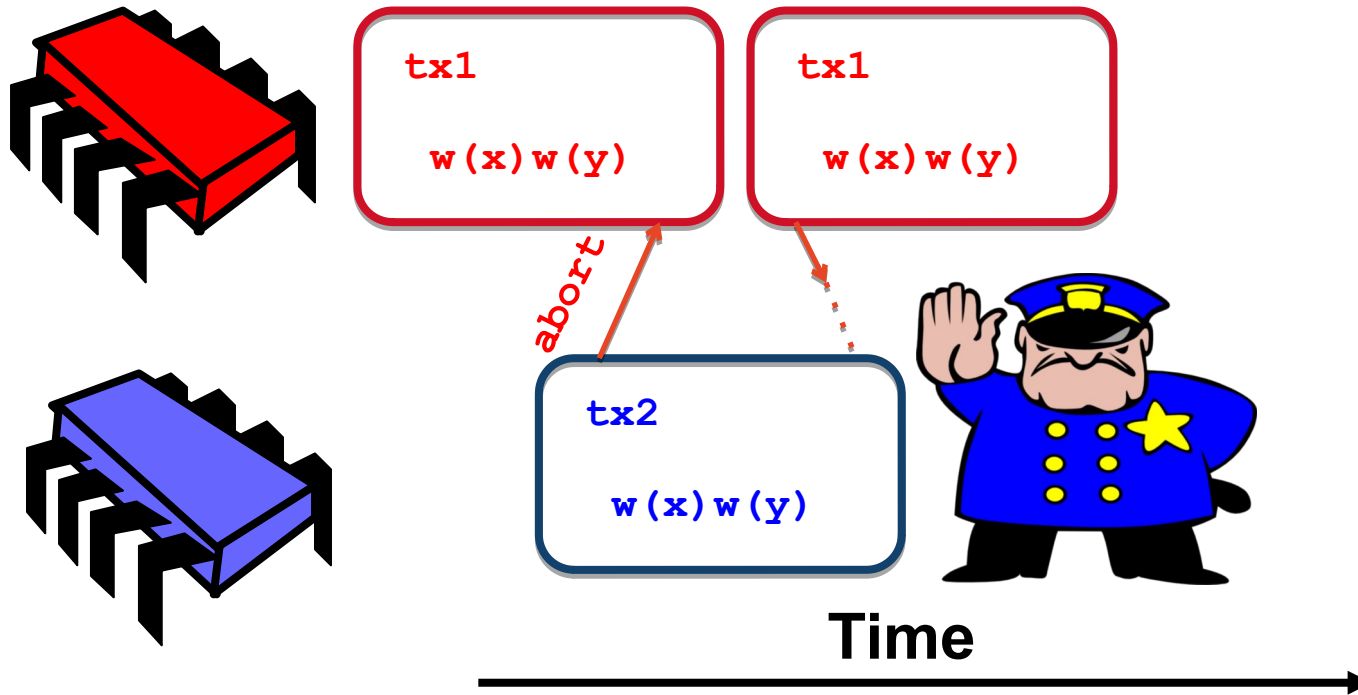

Starvation freedom

- Starvation: may happen if transactions keep aborting each other



Starvation freedom

- Contention manager: an external module called by the transactional memory to arbitrate between conflicting transactions.



Irrevocability

- An action is *irrevocable* if once it has been invoked it cannot be rolled back.
- As usual transactions execute speculatively, may abort and restart, they cannot execute irrevocable actions.

```
atomic {  
    ...  
    fire-missile() ;  
    ...  
}
```



- *Privatisation* ensures that a transaction is the only one to access data, thus guaranteeing that the transaction will commit. With privatisation, transactions can execute legacy code (and irrevocable actions) safely.

Expressiveness

Expressiveness limitations

- Transactions use a balanced open-close block whereas mutual exclusions (i.e., locks) can be “interleaved”.

```
atomic {  
    read(x)  
    read(y)  
    read(z)  
}
```

```
lock(x)  
    read(x)  
lock(y)  
unlock(x)  
    read(y)  
lock(z)  
unlock(y)  
    read(z)  
unlock(z)
```

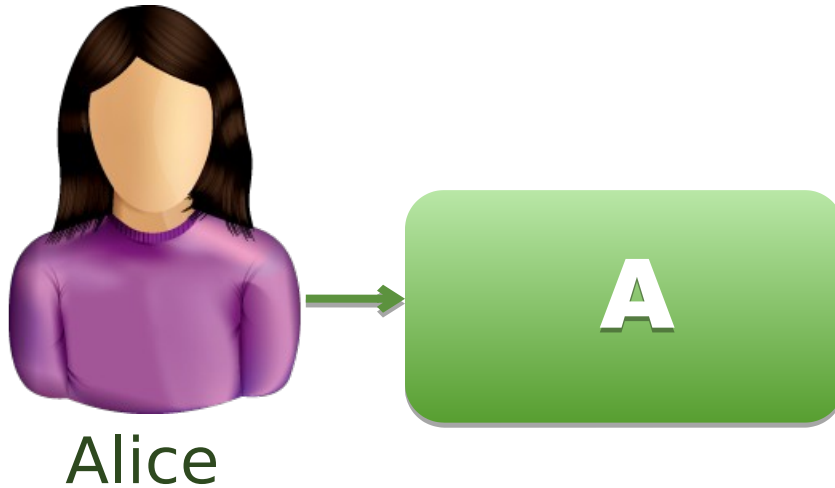
- *Relaxed transactions* help bridge the expressiveness gap between regular ones and locks.

Reusability

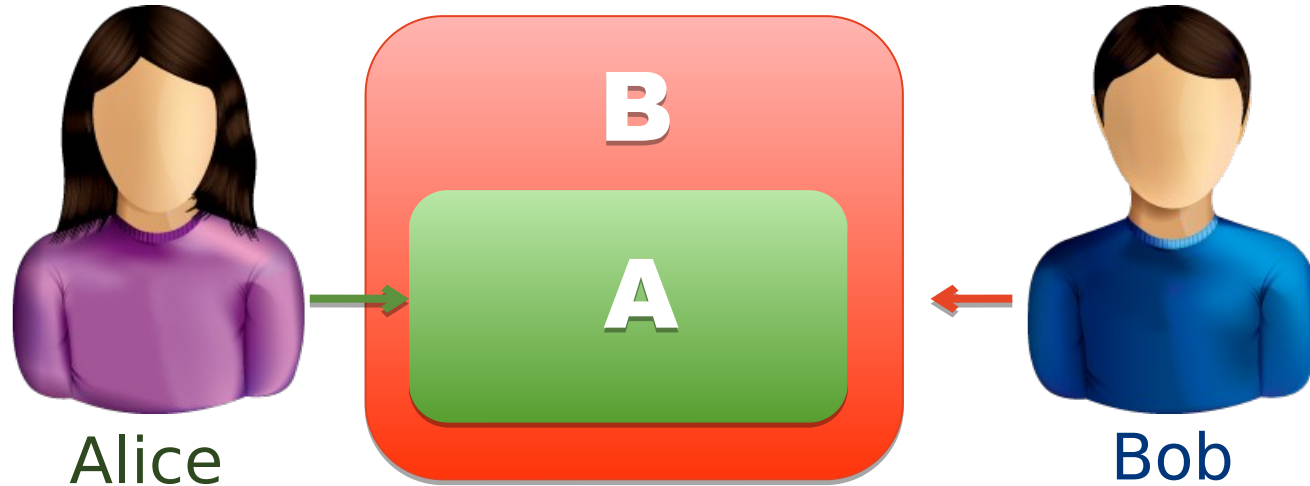


THE UNIVERSITY OF
SYDNEY

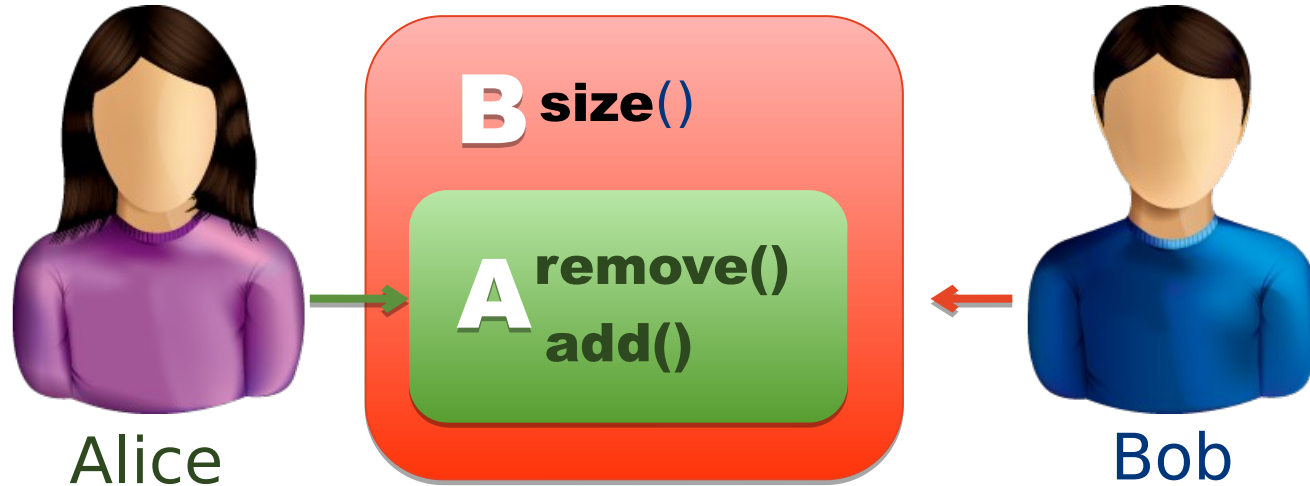
Reusability simplifies programming



Reusability simplifies programming

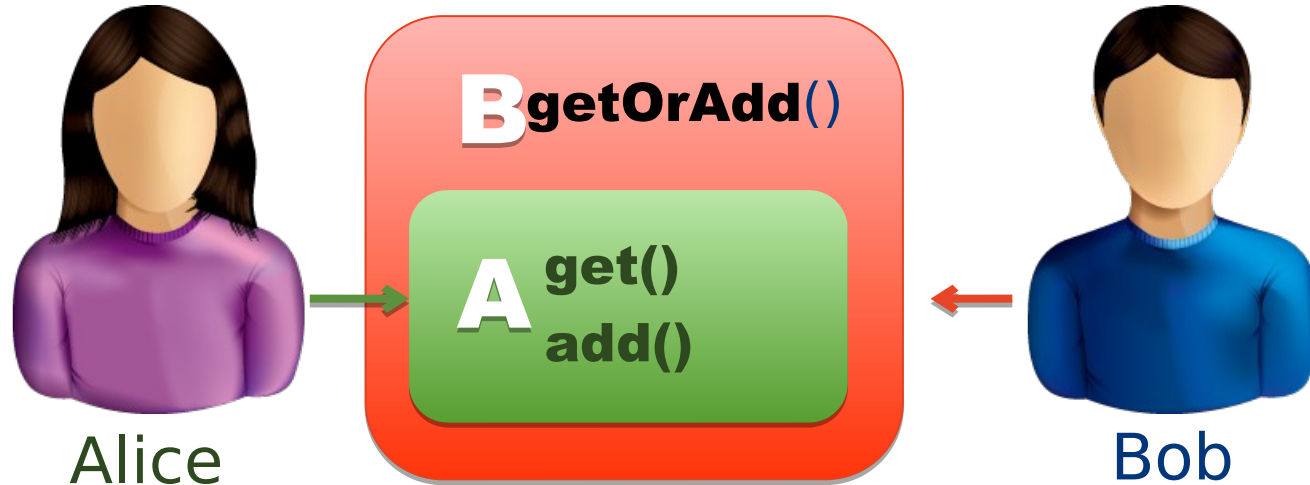


Reusability simplifies programming



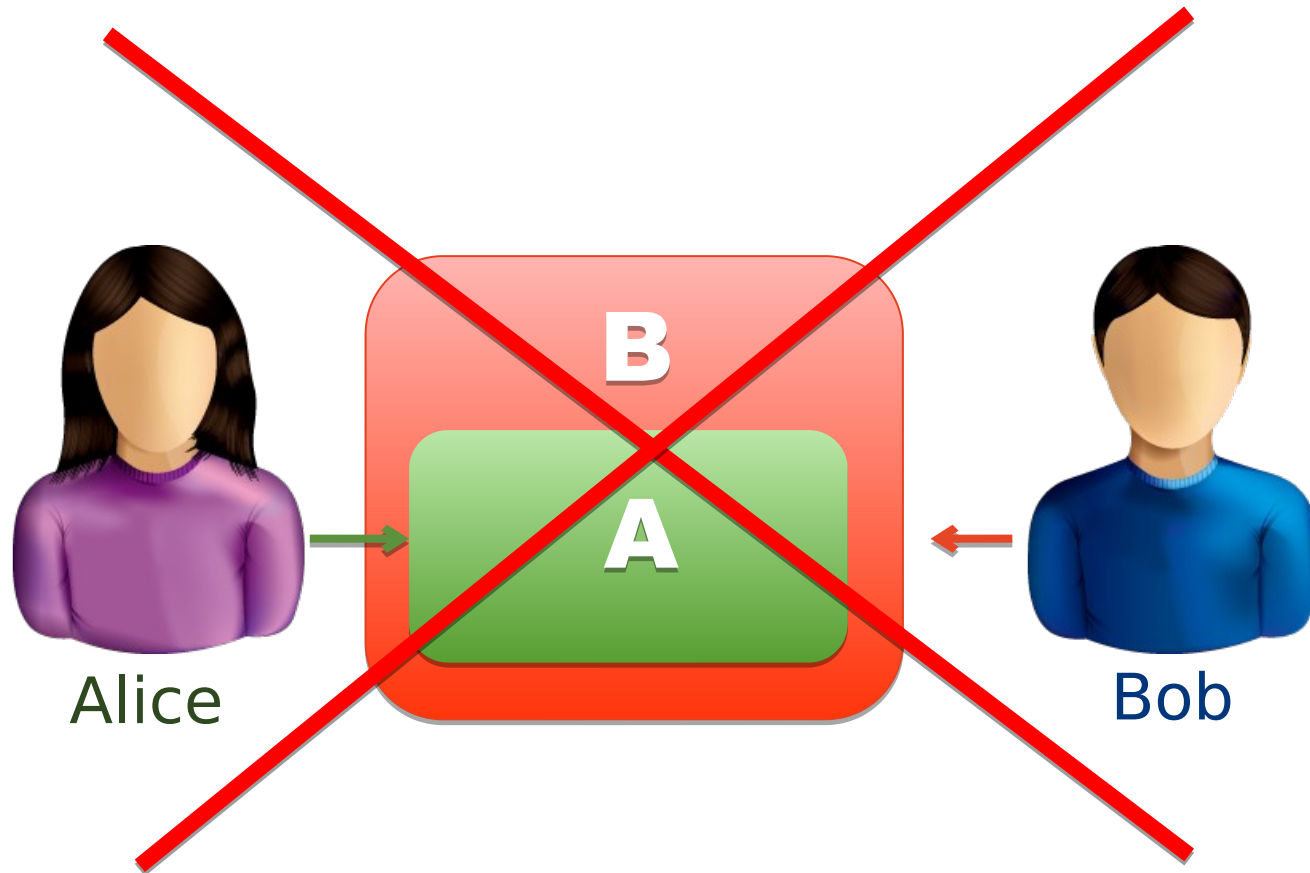
...through **extension** (adding a new method)

Reusability simplifies programming



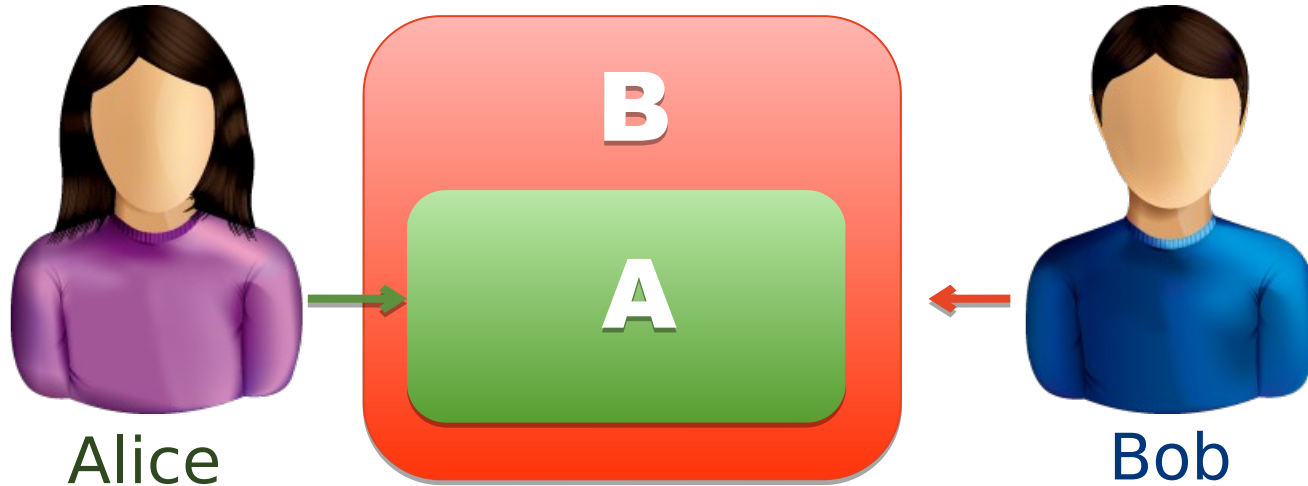
and through **composition** (composing methods into a new one)

Reusability simplifies programming



Not in a traditional concurrent environment

Reusability simplifies programming



Works with transactional memory

Software and Hardware TMs



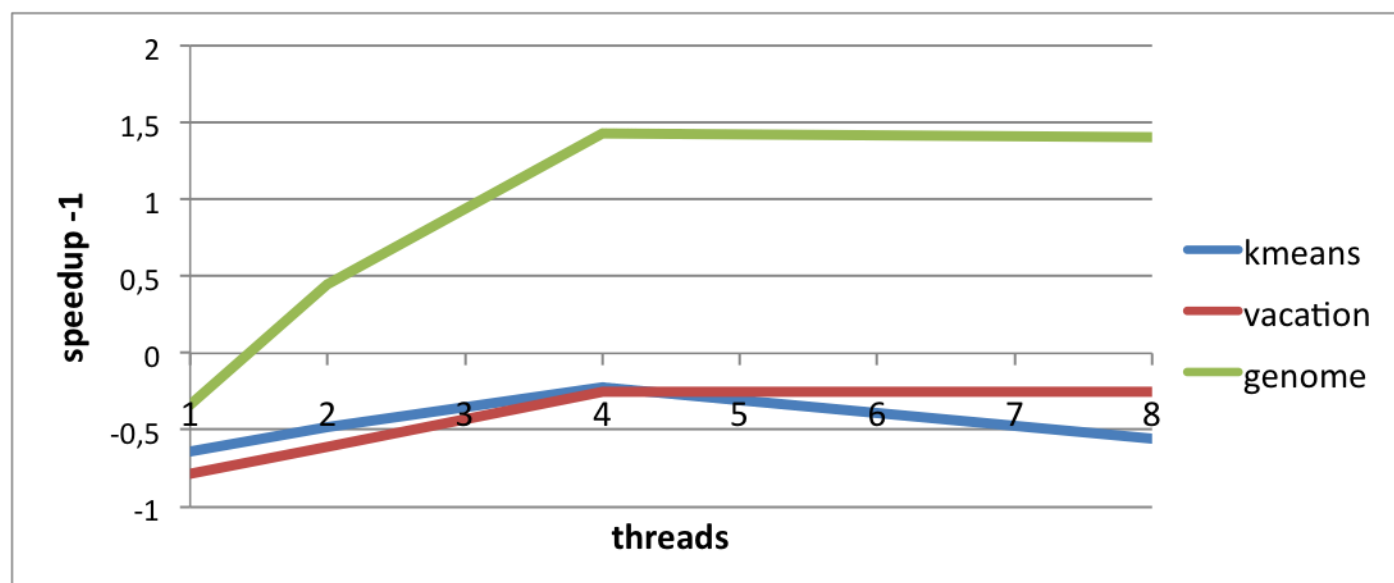
THE UNIVERSITY OF
SYDNEY

Overhead of Software TM (STM)

6 researchers from IBM claimed that “*STM is only a research toy*”
[IBM, CACM 2008] because:

Overhead of Software TM (STM)

6 researchers from IBM claimed that “*STM is only a research toy*” [IBM, CACM 2008] because:



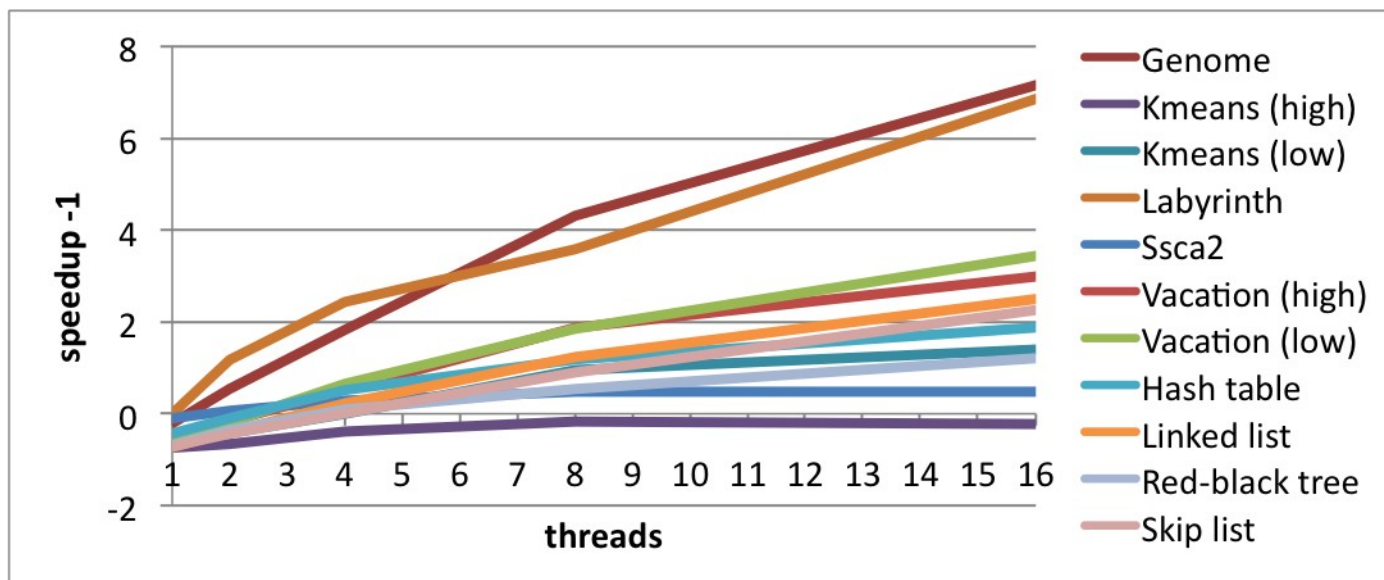
- STM overhead was too high to speed up sequential performance
- Performance was not even scaling to 8 threads

Overhead of Software TM (STM)

- We did a *thorough* analysis (benchmarks, STMs, architectures, compiler, privatisation, improvements)

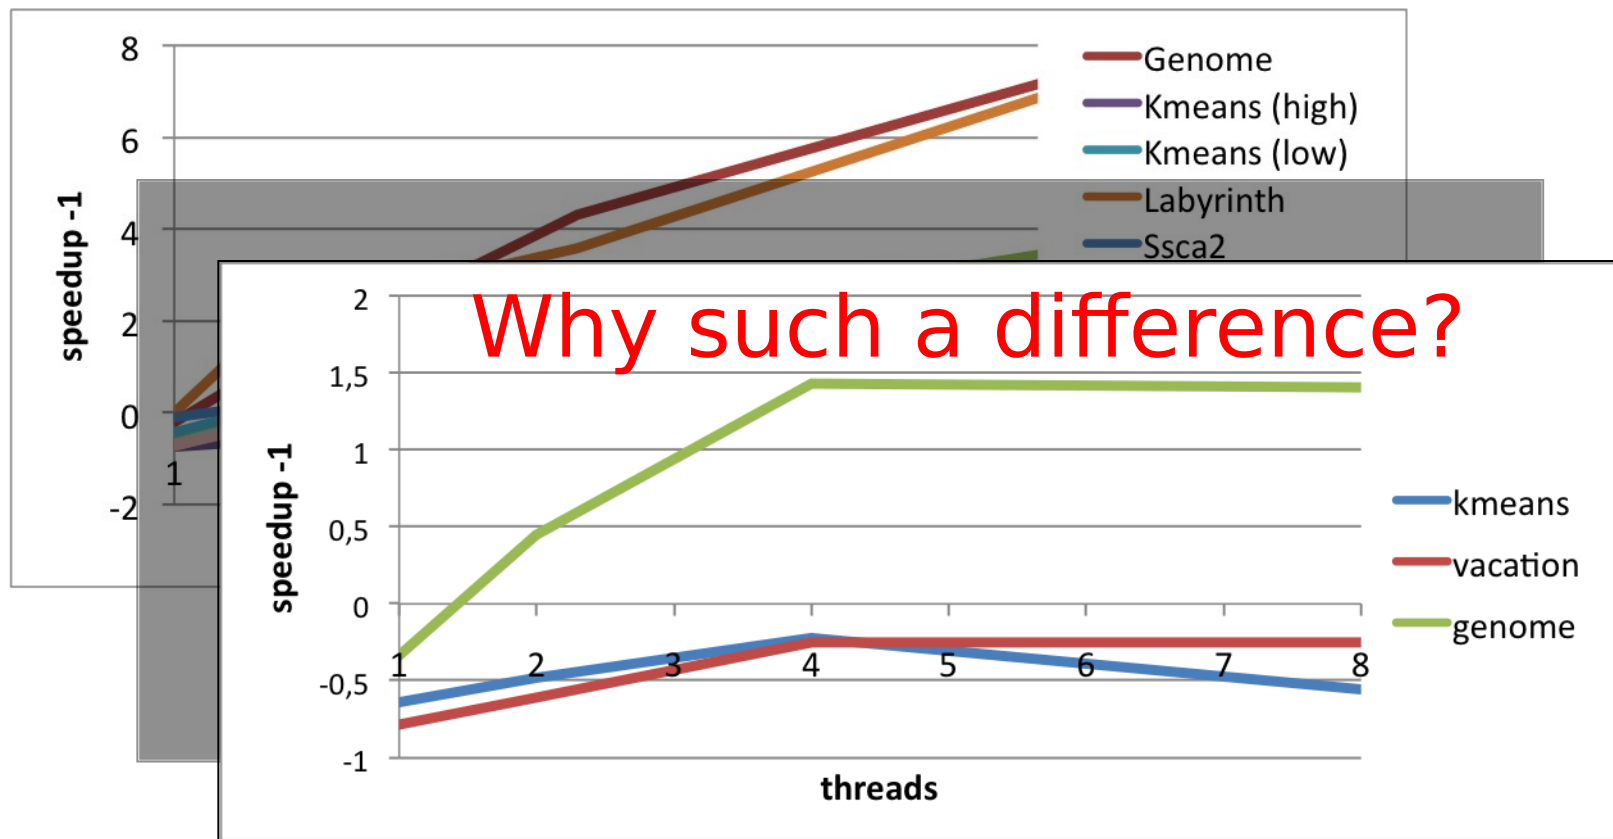
Overhead of Software TM (STM)

- We did a *thorough* analysis (benchmarks, STMs, architectures, compiler, privatization, improvements)



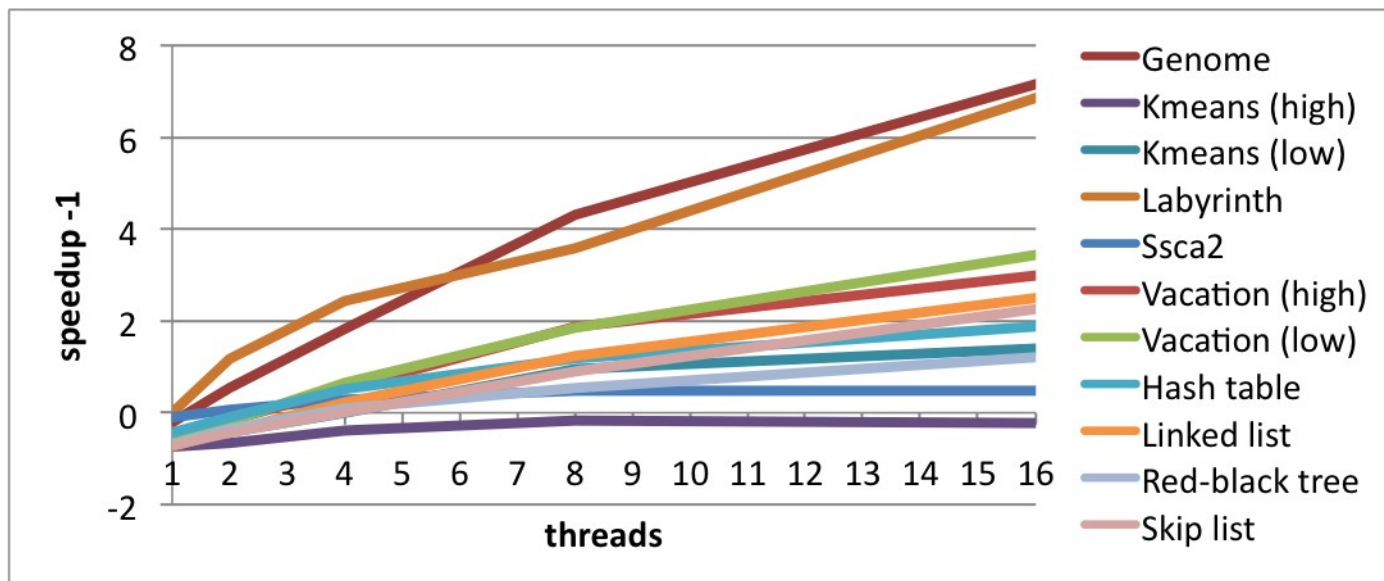
Overhead of Software TM (STM)

- We did a *thorough* analysis (benchmarks, STMs, architectures, compiler, privatization, improvements)



Overhead of Software TM (STM)

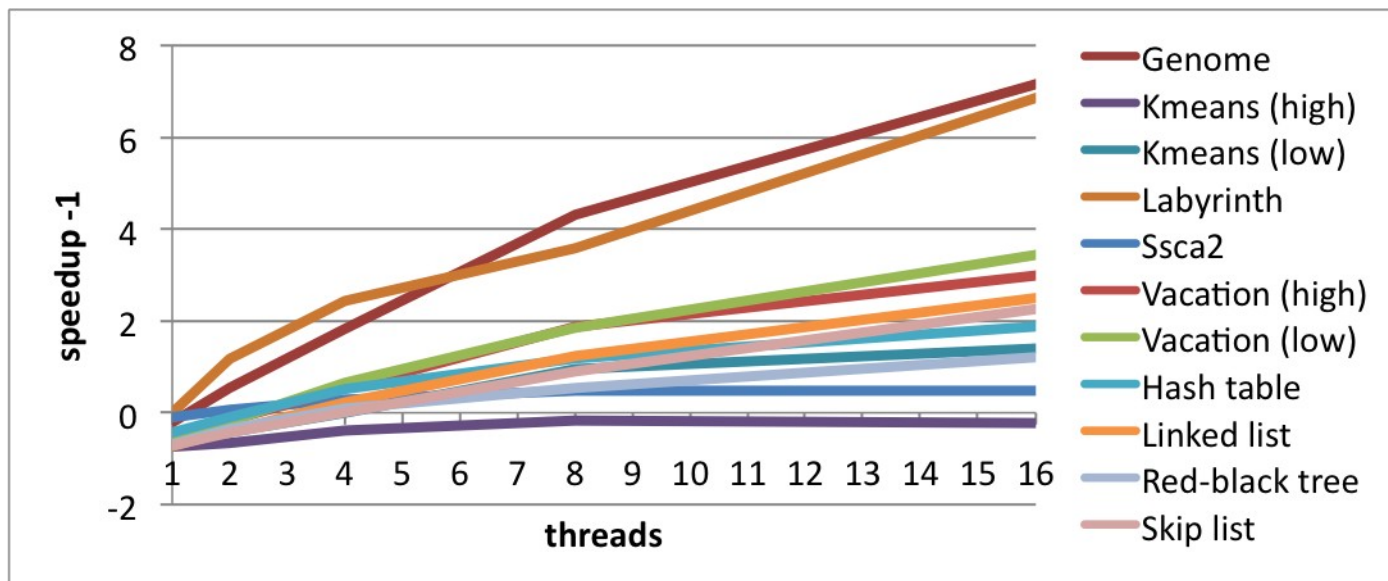
- We did a *thorough* analysis (benchmarks, STMs, architectures, compiler, privatization, improvements)



- Problem: 4 hyperthreaded cores \neq 8 cores

Overhead of Software TM (STM)

- We did a *thorough* analysis (benchmarks, STMs, architectures, compiler, privatization, improvements)



- Problem: 4 hyperthreaded cores \neq 8 cores
- Our rebuttal, “*Why STM can be more than a research toy*”, was published [CACM '11]

Hardware TM (HTM)

- February 2012: Transactional Synchronization Extensions (TSX-NI).
 - Hardware Lock Elision (HLE)
 - Restricted Transactional Memory (RTM)
- August 2014: Intel announced a bug in the TSX implementation on current steppings of Haswell, Haswell-E, Haswell-EP and early Broadwell CPUs
- November 2014: Bug fixed, in the vPro-enabled Core M-5Y70 Broadwell CPU in November 2014.
- 2015: IBM Power8 supports Hardware Transactional Memory
- 2015: Intel Skylake has TSX re-enabled.

Why should we care?

- It should be easy to make concurrent programs **safe** and **live**:
 - Safe: it ensures atomicity of transactions
 - Live: it generally ensures that a transaction commits if executing in isolation for long enough and sometimes that every transaction eventually commits
- Reusability/Composability:
 - A transaction can encapsulate multiple ones so that no knowledge is a priori necessary
 - Which makes concurrent programming modular
- Empirical studies indicate that concurrent programming with TM is:
 - Simpler
 - Safer
 - Faster ...than with existing synchronisation techniques.