

INFO1105/1905

Data Structures

Week 9: Graphs (start)

see textbook section 14.1, 14.2, 14.3

Professor Alan Fekete


Professor Seokhee Hong

School of Information Technologies

using material from the textbook
and A/Prof Kalina Yacef



THE UNIVERSITY OF
SYDNEY



Copyright warning

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

WARNING

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (**the Act**).

The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice.

- These slides contain material from the textbook (Goodrich, Tamassia & Goldwasser)
 - Data structures and algorithms in Java (5th & 6th edition)
- With modifications and additions from the University of Sydney
- The slides are a guide or overview of some big ideas
 - Students are responsible for knowing what is in the referenced sections of the textbook, not just what is in the slides

Reminder! Quiz 4

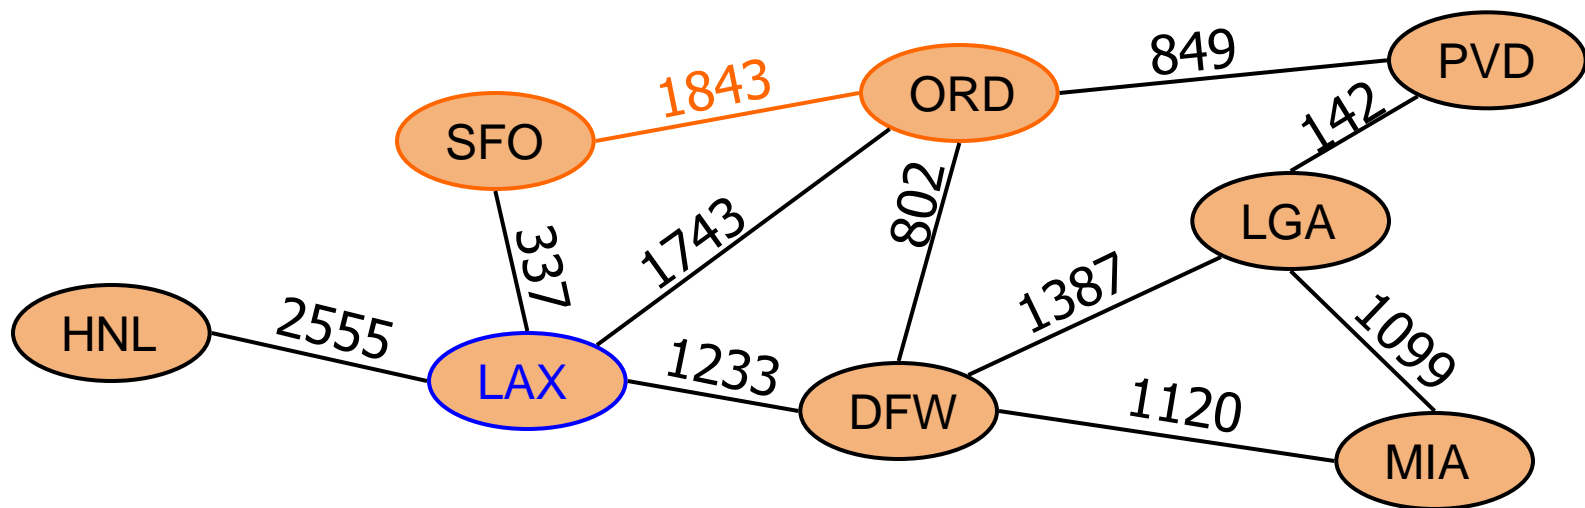
- Quiz 4 will take place during lab in week 10
- Done online, over a 20 minutes duration,
 - during the last 30 minutes of the lab period, or as indicated by your tutor
- A few multiple choice questions,
 - covering material from weeks 7, 8, also part of week 9
 - hash tables and collision handling (including chaining and open addressing)
 - recurrence equations
 - graph definition
 - graph representations
 - However, graph traversal (BFS, DFS) will be in quiz 5!

Outline

- **Graphs:**
 1. **Definitions**
 2. **Graph ADT**
- Data structures for graphs
- Graphs traversals


Graphs

- Graph: represent relationships between objects (vertices)
- A graph is a pair (V, E) , where
 - V is a set of nodes, called **vertices** (singular : **vertex**)
 - E is a collection of pairs of vertices, called **edges**
 - Vertices and edges are positions and store elements
- Example:
 - A **vertex** represents an airport and stores the three-letter airport code
 - An **edge** represents a flight route between two airports and stores the mileage of the route



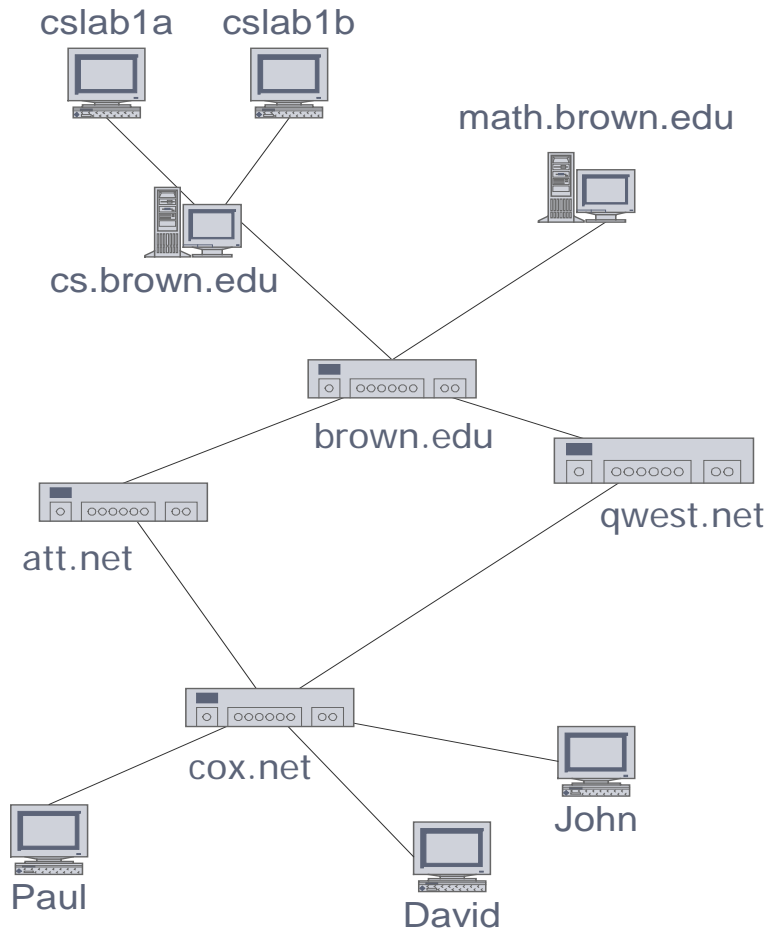
Edge Types

- **Directed edge**
 - ordered pair of vertices (u,v)
 - first vertex u is the **origin**
 - second vertex v is the **destination**
 - e.g., a flight
- **Undirected edge**
 - unordered pair of vertices (u,v)
 - e.g., a coauthorship relationship
- **Directed graph**
 - all the edges are directed
 - e.g., route network
- **Undirected graph**
 - all the edges are undirected
 - e.g., collaboration network

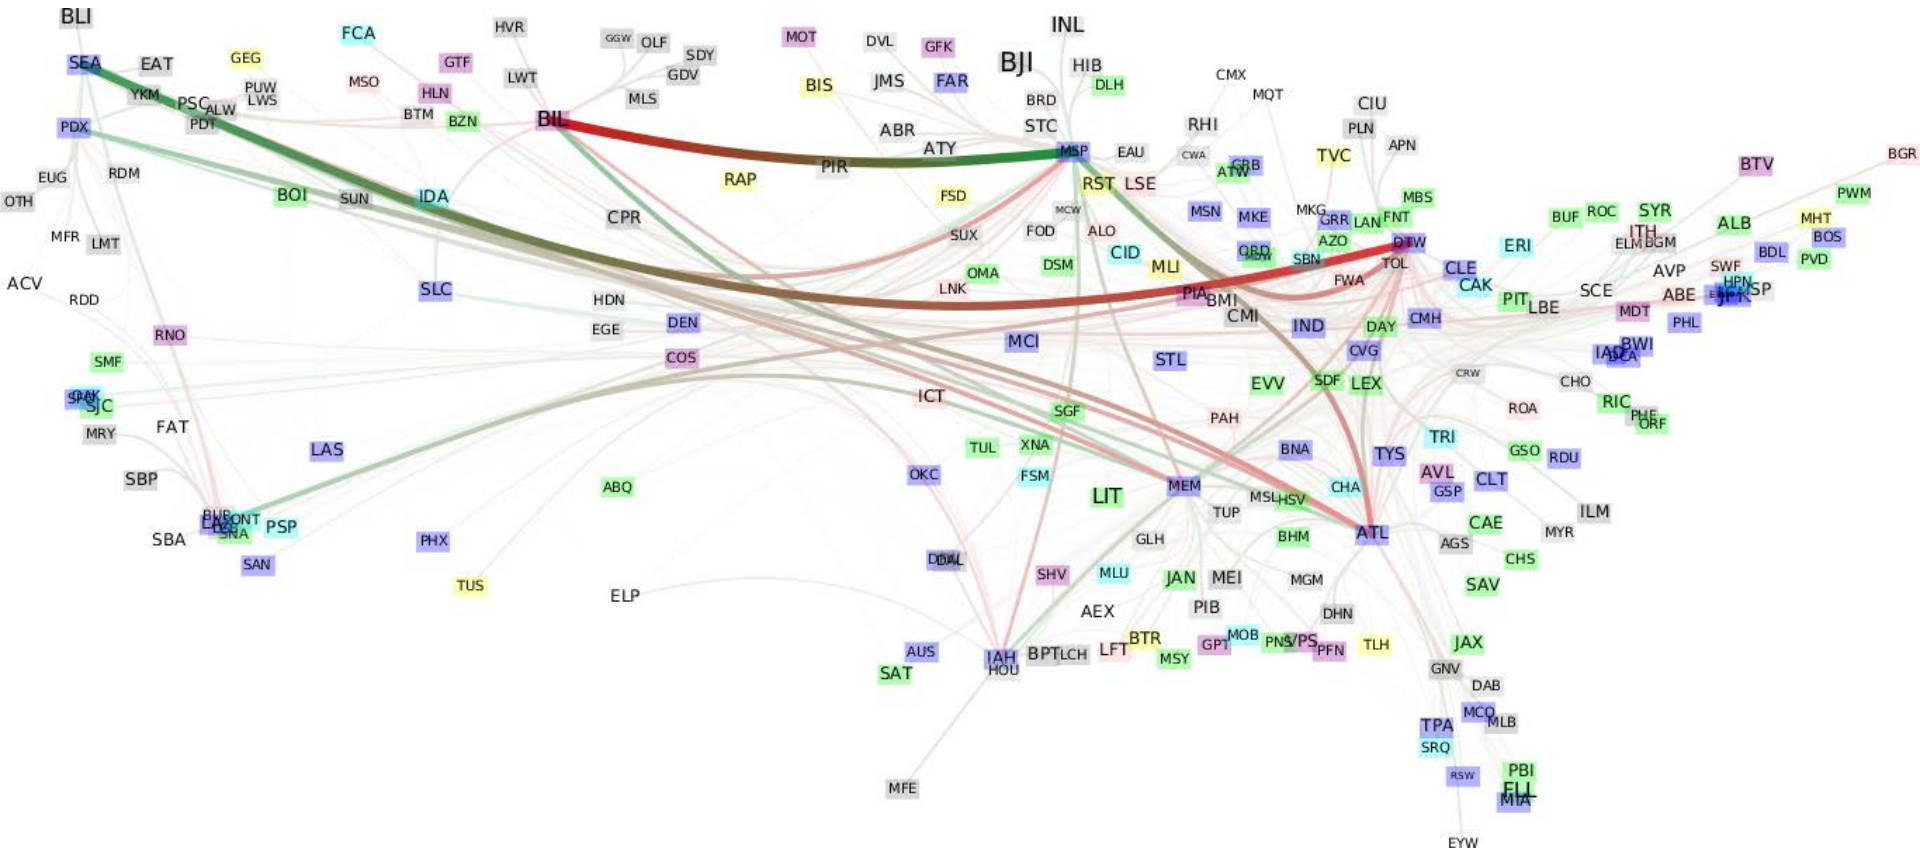


Applications

- Electronic circuits
 - Printed circuit board
 - Integrated circuit
- Transportation networks
 - Highway network
 - Flight network
 - City maps
- Computer networks
 - Local area network
 - Internet
 - Web
- Databases
 - Entity-relationship diagram
- OO programming
 - Class inheritance




US Airline Network



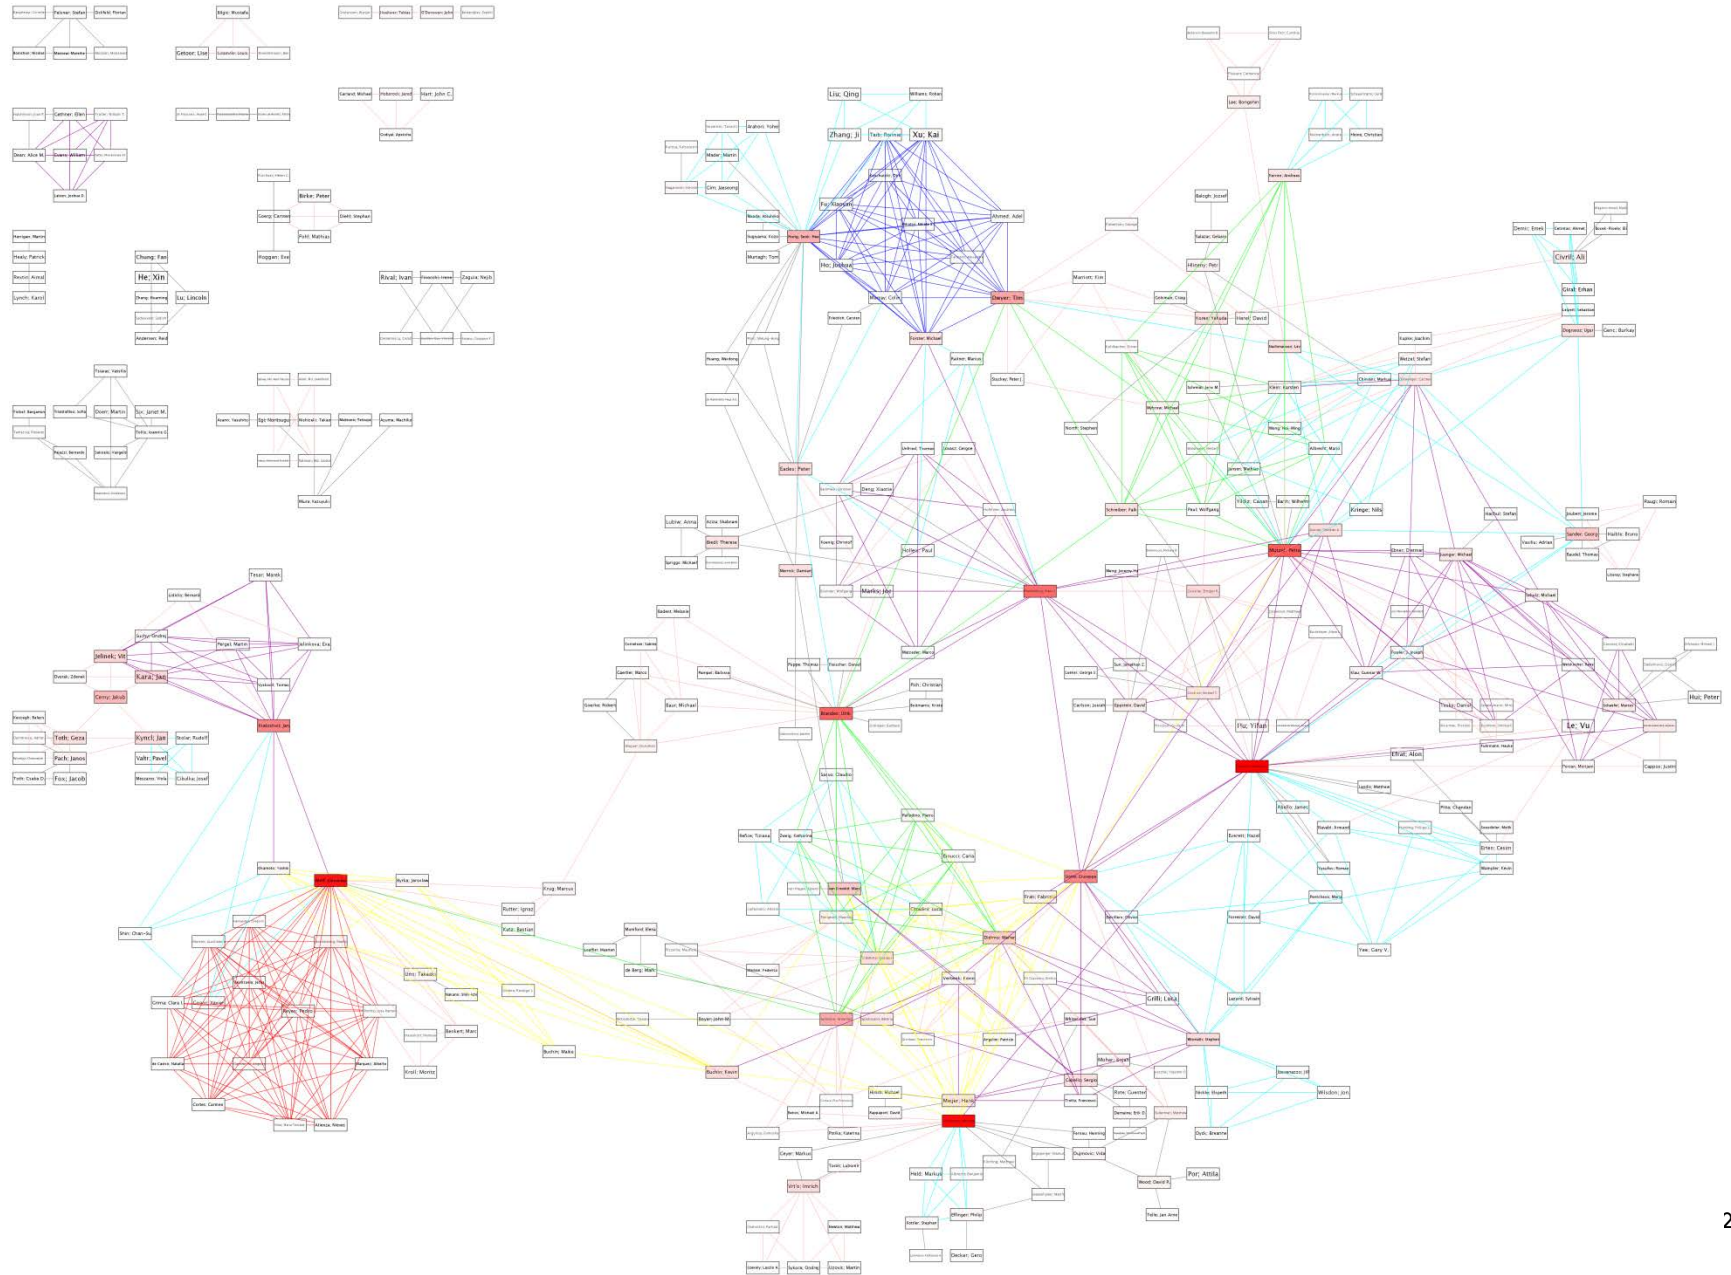
London Metro Map




Social Network: Facebook Network




Social Network: Research Collaboration Network



Biological Network: Protein-Protein Interaction Network




Internet Mapping Project



Terminology


- **End vertices** (or endpoints) of an edge
 - U and V are the endpoints of a
- Edges **incident** to a vertex
 - a, d, and b are incident on V
- **Adjacent** vertices
 - U and V are adjacent
- **Degree** of a vertex
 - X has degree 5
- **Parallel** edges
 - h and i are parallel edges
- **Self-loop**
 - j is a self-loop
- **Simple** graph
 - no parallel edges or self-loops



Terminology (cont.)

If edge is **directed**

- **Origin, destination** vertices
- **Outgoing** edges of V are a, d
- **Incoming** edge of V is b
- Degree of a vertex
 - $\deg(V)$ is 3
 - **indeg**(V) is 1: in-degree
 - **outdeg**(V) is 2: out-degree



Terminology (cont.)

– Path

- sequence of alternating vertices and edges
- begins with a vertex
- ends with a vertex
- each edge is preceded and followed by its endpoints


– Simple path

- path such that all its vertices and edges are distinct

– Examples

- $P_1 = (V, b, X, h, Z)$ is a simple path
- $P_2 = (U, c, W, e, X, g, Y, f, W, d, V)$ is a path that is not simple

- In directed graph, **directed** paths



Terminology (cont.)

– Cycle

- circular sequence of alternating vertices and edges
- Path that starts and ends at the same vertex


– Simple cycle

- cycle such that all its vertices and edges are distinct

– Examples

- $C_1 = (V, b, X, g, Y, f, W, c, U, a, \downarrow)$ is a simple cycle
- $C_2 = (U, c, W, e, X, g, Y, f, W, d, V, a, \downarrow)$ is a cycle that is not simple

– Acyclic graph has no cycle



Properties

Property 1

$$\sum_v \deg(v) = 2m$$

Proof: each edge is counted twice

Property 2

In an undirected simple graph

$$m \leq n(n-1)/2$$

Proof: each vertex has degree at most $(n-1)$

directed simple graph:


$$m \leq n(n-1)$$

* Complete graph K_n (clique):

$$m = n(n-1)/2$$

Notation

n	number of vertices
m	number of edges
$\deg(v)$	degree of vertex v



Example: K_4

- $n = 4$
- $m = 6$
- $\deg(v) = 3$

Graph ADT: Vertices and Edges

- A **graph** is a collection of **vertices** and **edges**.
- We model the abstraction as a combination of three data types: **Vertex**, **Edge**, and **Graph**.
- A **Vertex** is a lightweight object that stores an arbitrary element provided by the user (e.g., an airport code)
 - We assume it supports a method, `getElement()`, to retrieve the stored element.
- An **Edge** stores an associated object (e.g., a flight number, travel distance, cost), retrieved with the `getElement()` method.

Graph ADT Methods

- Old version

`areAdjacent(v, w):`
`getEdge(u, v)`

`incidentEdges(v):`
`outgoingEdges(v)`
`incomingEdges(v)`



`numVertices():` Returns the number of vertices of the graph.
`vertices():` Returns an iteration of all the vertices of the graph.
`numEdges():` Returns the number of edges of the graph.
`edges():` Returns an iteration of all the edges of the graph.
`getEdge(u, v):` Returns the edge from vertex u to vertex v , if one exists; otherwise return null. For an undirected graph, there is no difference between `getEdge(u, v)` and `getEdge(v, u)`.
`endVertices(e):` Returns an array containing the two endpoint vertices of edge e . If the graph is directed, the first vertex is the origin and the second is the destination.
`opposite(v, e):` For edge e incident to vertex v , returns the other vertex of the edge; an error occurs if e is not incident to v .
`outDegree(v):` Returns the number of outgoing edges from vertex v .
`inDegree(v):` Returns the number of incoming edges to vertex v . For an undirected graph, this returns the same value as does `outDegree(v)`.
`outgoingEdges(v):` Returns an iteration of all outgoing edges from vertex v .
`incomingEdges(v):` Returns an iteration of all incoming edges to vertex v . For an undirected graph, this returns the same collection as does `outgoingEdges(v)`.
`insertVertex(x):` Creates and returns a new Vertex storing element x .
`insertEdge(u, v, x):` Creates and returns a new Edge from vertex u to vertex v , storing element x ; an error occurs if there already exists an edge from u to v .
`removeVertex(v):` Removes vertex v and all its incident edges from the graph.
`removeEdge(e):` Removes edge e from the graph.

Outline

- Graphs: definitions and ADT
- **Data structures for graphs**
 1. **Edge list structure**
 2. **Adjacency list structure**
 3. **Adjacency map structure**
 4. **Adjacency matrix**
- Graphs traversals



1. Edge List Structure

- Simple, but inefficient
- **Unordered list V, E**
- **Vertex object**
 - element
 - reference to position in vertex sequence
- **Edge object**
 - element
 - **origin** vertex object
 - **destination** vertex object
 - reference to position in edge sequence
- **Vertex sequence V**
 - sequence of vertex objects
- **Edge sequence E**
 - sequence of edge objects



Edge List Structure

- ▶ Unordered list of all edges
- ▶ Vertex object
 - ▶ element
 - ▶ reference to position in vertex sequence
- ▶ Edge object
 - ▶ element
 - ▶ origin vertex object
 - ▶ destination vertex object
 - ▶ reference to position in edge sequence
- ▶ Vertex sequence V
 - ▶ sequence of vertex objects
- ▶ Edge sequence E
 - ▶ sequence of edge objects



Edge list: performance

Method	Edge List
Space	$n + m$
<code>incidentEdges(v)</code>	m
<code>areAdjacent(v, w)</code>	m
<code>insertVertex(o)</code>	1
<code>insertEdge(v, w, o)</code>	1
<code>removeVertex(v)</code>	m
<code>removeEdge(e)</code>	1

- n vertices, m edges
- no parallel edges
- no self-loops



* Method

`areAdjacent(v, w):`
`getEdge(u, v)`

`incidentEdges(v):`
`outgoingEdges(v)`
`incomingEdges(v)`


2. Adjacency List Structure

- Unordered list with additional list structure
- Incidence sequence for each vertex v
 - sequence of references to edge objects of incident edges
- Augmented edge objects
 - references to associated positions in incidence sequences of end vertices
- Sparse graphs: $O(n)$ edges



Adjacency List Structure

- Unordered list with additional list structure
- Incidence sequence $I(v)$ for each vertex v
 - sequence of references to edge objects of incident edges
- Augmented edge objects
 - references to associated positions in incidence sequences of end vertices



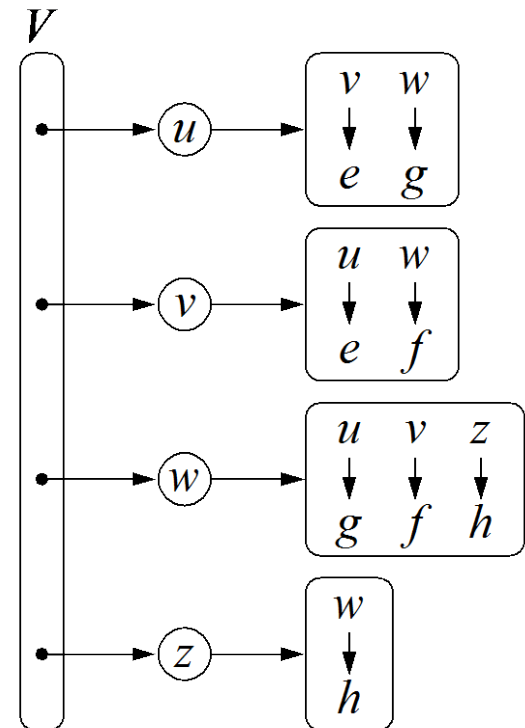

Adjacency list: performance

- All methods in $O(1)$ with Edge list structure still $O(1)$

Method	Adjacency List
Space	$n + m$
<code>incidentEdges(v)</code>	$\text{deg}(v)$
<code>areAdjacent (v, w)</code>	$\min(\text{deg}(v), \text{deg}(w))$
<code>insertVertex(o)</code>	1
<code>insertEdge(v, w, o)</code>	1
<code>removeVertex(v)</code>	$\text{deg}(v)$
<code>removeEdge(e)</code>	1


3. Adjacency Map Structure

- ▶ Same as adjacency list, but uses a hash-based map for storing incident edges (see week 8)
- ▶ Incidence map for each vertex v
 - ▶ Key = opposite endpoint Value = edge
- ▶ getEdge(u,v) now is in expected $O(1)$
 - ▶ Although still worst case $O(\min(\deg(u), \deg(v)))$



4. Adjacency Matrix Structure

- Augmented vertex objects
 - Integer key (index) associated with vertex
- 2D-array adjacency array A
 - Reference to edge object for adjacent vertices
 - Null for non nonadjacent vertices
- The “old fashioned” version just has 0 for no edge and 1 for edge
- Dense graphs: $O(n^2)$ edges



		0	1	2	3		
u	\longrightarrow	0		e	g		
v	\longrightarrow	1		e		f	
w	\longrightarrow	2		g	f		h
z	\longrightarrow	3				h	

Adjacency Matrix: performance

Method	Adjacency Matrix
Space	n^2
<code>incidentEdges(<i>v</i>)</code>	n
<code>areAdjacent (<i>v</i>, <i>w</i>)</code>	1
<code>insertVertex(<i>o</i>)</code>	n^2
<code>insertEdge(<i>v</i>, <i>w</i>, <i>o</i>)</code>	1
<code>removeVertex(<i>v</i>)</code>	n^2
<code>removeEdge(<i>e</i>)</code>	1

Performance: summary

- n vertices, m edges, no parallel edges/ self-loops
- `incidentEdges(v)`: `outgoingEdges(v)`, `incomingEdges(v)`


Method	Edge List	Adjacency List	Adjacency Map	Adjacency Matrix
Space	$n + m$	$n + m$	$n + m$	n^2
<code>getEdge(u, v)</code>	m	$\min(\deg(v), \deg(w))$	1 (exp.)	1
<code>outDegree(v), inDegree(v)</code>	m	1	1	n
<code>incidentEdges(v)</code>	m	$\deg(v)$	$\deg(v)$	n
<code>areAdjacent(v, w)</code>	m	$\min(\deg(v), \deg(w))$	1 (exp.)	1
<code>insertVertex(o)</code>	1	1	1	n^2
<code>insertEdge(v, w, o)</code>	1	1	1 (exp.)	1
<code>removeVertex(v)</code>	m	$\deg(v)$	$\deg(v)$	n^2
<code>removeEdge(e)</code>	1	1	1	1

Outline


- Graphs: definitions and ADT
- Data structures for graphs
- **Graphs traversals**
 1. **More definitions**
 2. **DFS**
 3. **BFS**

Subgraphs

- A **subgraph** S of a graph G is a graph such that
 - The vertices of S are a subset of the vertices of G
 - The edges of S are a subset of the edges of G
- A **spanning subgraph** of G is a subgraph that contains all the vertices of G




Subgraph




Spanning subgraph

Connectivity

- A graph is **connected** if there is a path between every pair of vertices
- A **connected component** of a graph G is a *maximal* connected subgraph of G



Connected graph




Non connected graph with two connected components

Trees and Forests


- A **(free) tree** is an undirected graph T such that
 - T is connected
 - T has no cycles

This definition of tree is different from the one of a **rooted tree**



Tree


- A **forest** is an undirected graph without cycles
- The connected components of a forest are trees




Forest

Spanning Trees and Forests

- A **spanning tree** of a connected graph is a spanning subgraph that is a tree
 - A spanning tree is not unique unless the graph is a tree
 - Spanning trees have applications to the design of communication networks
 - A **spanning forest** of a graph is a spanning subgraph that is a forest
-
- Properties for G (undirected) graph with n vertices and m edges
 - If G is connected then $m \geq n-1$
 - If G is a (free) tree then $m = n-1$
 - If G is a forest then $m \leq n-1$




Graph



Spanning tree

Graph Traversal: DFS, BFS

- Graph traversal algorithms are key to answering “Reachability” problems:
 - How to travel from one vertex to another following paths of a graph
- For example: given graph G (undirected/directed)
 - Find and report a path between two given vertices
 - Compute a shortest path (min # of edges) between two vertices
 - Test whether G is connected/strongly-connected
 - Computes a spanning tree/forest of G
 - Computes the (strongly) connected components of G
 - Find a cycle in the graph



Depth-First Search

Depth-First Search

- Depth-first search (DFS) is a general technique for traversing a graph
- A DFS traversal of a graph G
 - Visits all the vertices and edges of G
- DFS on a graph with n vertices and m edges takes $O(n + m)$ time
- DFS can be further extended to solve other graph problems
 - Find and report a path between two given vertices
 - Find a cycle in the graph
 - Determines whether G is connected
 - Computes the connected components of G
 - Computes a spanning forest of G

DFS Algorithm from a starting vertex u

Algorithm DFS(G, u):

Input: A graph G and a vertex u of G

Output: A collection of vertices reachable from u , with their discovery edges

Mark vertex u as visited.

for each of u 's outgoing edges, $e = (u, v)$ **do**

if vertex v has not been visited **then**






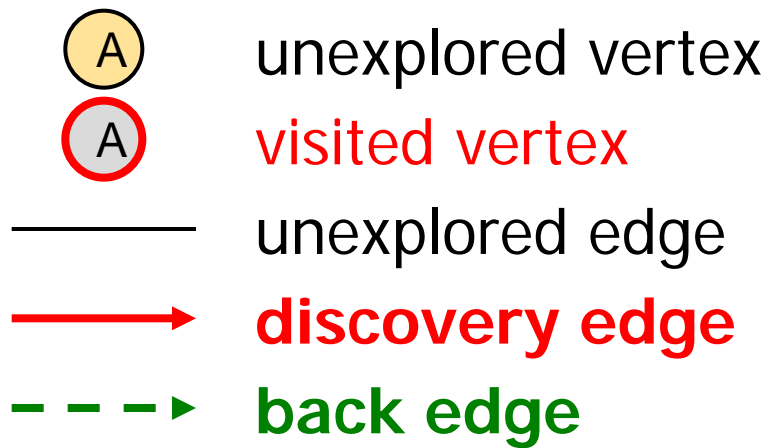
 Record edge e as the discovery edge for vertex v .

 Recursively call DFS(G, v).





Java implementation (fragment 14.5)

```
1  /** Performs depth-first search of Graph g starting at Vertex u. */
2  public static <V,E> void DFS(Graph<V,E> g, Vertex<V> u,
3      Set<Vertex<V>> known, Map<Vertex<V>,Edge<E>> forest) {
4      known.add(u);                // u has been discovered
5      for (Edge<E> e : g.outgoingEdges(u)) { // for every outgoing edge from u
6          Vertex<V> v = g.opposite(u, e);
7          if (!known.contains(v)) {
8              forest.put(v, e);        // e is the tree edge that discovered v
9              DFS(g, v, known, forest); // recursively explore from v
10         }
11     }
12 }
```


Example




Example (cont.)



DFS and Maze Traversal




- The DFS algorithm is similar to a classic strategy for exploring a maze
 - We mark each intersection, corner and dead end (vertex) visited
 - We mark each corridor (edge) traversed
 - We keep track of the path back to the entrance (start vertex) by means of a rope (recursion stack)




DFS and Maze Traversal

- The DFS algorithm is similar to a classic strategy for exploring a maze
 - We mark each intersection, corner and dead end (vertex) visited
 - We mark each corridor (edge) traversed
 - We keep track of the path back to the entrance (start vertex) by means of a rope (recursion stack)




DFS and Maze Traversal

- The DFS algorithm is similar to a classic strategy for exploring a maze
 - We mark each intersection, corner and dead end (vertex) visited
 - We mark each corridor (edge) traversed
 - We keep track of the path back to the entrance (start vertex) by means of a rope (recursion stack)



DFS and Maze Traversal

- The DFS algorithm is similar to a classic strategy for exploring a maze
 - We mark each intersection, corner and dead end (vertex) visited
 - We mark each corridor (edge) traversed
 - We keep track of the path back to the entrance (start vertex) by means of a rope (recursion stack)



Properties of DFS

Property 1


$DFS(G, v)$ visits all the vertices and edges in the **connected component** of v

Property 2

The **discovery edges** labeled by $DFS(G, v)$ form a **spanning tree (DFS tree)** of the connected component of v

* Directed graphs: DFS tree T

1. Tree edges: discovery edges
2. Non-tree edges
 - **Back** edges: to ancestor in T
 - **Forward** edges: to descendant in T
 - **Cross** edges



Analysis of DFS

- Setting/getting a vertex/edge label takes $O(1)$ time
- DFS is called at most once for each vertex
- Each edge is examined at most twice (once for directed graph)
- Each vertex is labeled twice
 - once as UNEXPLORED
 - once as VISITED
- Each edge is labeled twice
 - once as UNEXPLORED
 - once as DISCOVERY or BACK
- Method incidentEdges (outgoingEdges(v)) takes $O(\deg(v))$ time
- DFS runs in $O(n + m)$ time provided the graph is represented by the adjacency list structure
 - Recall that $\sum_v \deg(v) = 2m$

Algorithm *DFS*(G, v)

Input graph G and a start vertex v of G

Output labeling of the edges of G
in the connected component of v
as discovery edges and back edges

setLabel(v , VISITED)

for all $e \in G.\text{incidentEdges}(v)$

if *getLabel*(e) = UNEXPLORED

$w \leftarrow \text{opposite}(v, e)$

if *getLabel*(w) = UNEXPLORED

setLabel(e , DISCOVERY)

DFS(G, w)

else

setLabel(e , BACK)

DFS Application

- DFS can be further extended to solve other graph problems in $O(n + m)$ time
 - Find and report a **path** between two given vertices
 - Find a **cycle** in the graph
 - Test whether G is **connected**
 - Computes the **connected components** of G
 - Computes a **spanning tree** of G (if G is connected)
- DFS Implementation (see DFS method in Section 14.3.2)
 - Keep track of DFS tree: two auxiliary data structures
 - a **set** (**known**): vertices that have already been visited
 - a **map** (**forest**) : associates a vertex v and edge e (used to discover v)
 - Hash-based implementation: $O(1)$ expected time (mark vertex v “explored”)

1. Path Finding

- We can specialize the DFS algorithm to find a path between two given vertices u and z using the template method pattern
- We call $DFS(G, u)$ with u as the start vertex
- We use a stack S to keep track of the path between the start vertex and the current vertex
- As soon as destination vertex z is encountered, we return the path as the contents of the stack

```
Algorithm pathDFS( $G, v, z$ )  
  setLabel( $v, VISITED$ )  
  S.push( $v$ )  
  if  $v = z$   
    return S.elements()  
  for all  $e \in G.incidentEdges(v)$   
    if getLabel( $e$ ) = UNEXPLORED  
       $w \leftarrow opposite(v, e)$   
      if getLabel( $w$ ) = UNEXPLORED  
        setLabel( $e, DISCOVERY$ )  
        S.push( $e$ )  
        pathDFS( $G, w, z$ )  
        S.pop( $e$ )  
      else  
        setLabel( $e, BACK$ )  
  S.pop( $v$ )
```

Path Finding in Java

```
1  /** Returns an ordered list of edges comprising the directed path from u to v. */
2  public static <V,E> PositionalList<Edge<E>>
3  constructPath(Graph<V,E> g, Vertex<V> u, Vertex<V> v,
4               Map<Vertex<V>,Edge<E>> forest) {
5      PositionalList<Edge<E>> path = new LinkedPositionalList<>();
6      if (forest.get(v) != null) {           // v was discovered during the search
7          Vertex<V> walk = v;                // we construct the path from back to front
8          while (walk != u) {
9              Edge<E> edge = forest.get(walk);
10             path.addFirst(edge);            // add edge to *front* of path
11             walk = g.opposite(walk, edge);  // repeat with opposite endpoint
12         }
13     }
14     return path;
15 }
```

2. Cycle Finding

- We can specialize the DFS algorithm to find a simple cycle using the template method pattern
- We use a stack S to keep track of the path between the start vertex and the current vertex
- A cycle exists iff a back edge exists
- As soon as a back edge (v, w) is encountered, we return the cycle as the portion of the stack from the top to vertex w

```
Algorithm cycleDFS( $G, v, z$ )  
  setLabel( $v, VISITED$ )  
   $S.push(v)$   
  for all  $e \in G.incidentEdges(v)$   
    if getLabel( $e$ ) = UNEXPLORED  
       $w \leftarrow opposite(v, e)$   
       $S.push(e)$   
      if getLabel( $w$ ) = UNEXPLORED  
        setLabel( $e, DISCOVERY$ )  
        pathDFS( $G, w, z$ )  
         $S.pop(e)$   
      else  
         $T \leftarrow$  new empty stack  
        repeat  
           $o \leftarrow S.pop()$   
           $T.push(o)$   
        until  $o = w$   
        return  $T.elements()$   
   $S.pop(v)$ 
```

DFS for an Entire Graph

- The algorithm uses a mechanism for setting and getting “labels” of vertices and edges

Algorithm *DFS(G)*

Input graph *G*

Output labeling of the edges of *G*
as discovery edges and
back edges

```
for all u ∈ G.vertices()
    setLabel(u, UNEXPLORED)
for all e ∈ G.edges()
    setLabel(e, UNEXPLORED)
for all v ∈ G.vertices()
    if getLabel(v) = UNEXPLORED
        DFS(G, v)
```

Algorithm *DFS(G, v)*


Input graph *G* and a start vertex *v* of *G*
Output labeling of the edges of *G*
in the connected component of *v*
as discovery edges and back edges

```
setLabel(v, VISITED)
for all e ∈ G.incidentEdges(v)
    if getLabel(e) = UNEXPLORED
        w ← opposite(v, e)
        if getLabel(w) = UNEXPLORED
            setLabel(e, DISCOVERY)
            DFS(G, w)
        else
            setLabel(e, BACK)
```

3. All Connected Components

- Loop over all vertices, doing a DFS from each unvisited one.

```
1  /** Performs DFS for the entire graph and returns the DFS forest as a map. */
2  public static <V,E> Map<Vertex<V>,Edge<E>> DFSComplete(Graph<V,E> g) {
3      Set<Vertex<V>> known = new HashSet<>();
4      Map<Vertex<V>,Edge<E>> forest = new ProbeHashMap<>();
5      for (Vertex<V> u : g.vertices())
6          if (!known.contains(u))
7              DFS(g, u, known, forest);           // (re)start the DFS process at u
8      return forest;
9  }
```

Breadth-First Search

Breadth-First Search

- Breadth-first search (BFS) is a general technique for traversing a graph
- A BFS traversal of a graph G
 - Visits all the vertices and edges of G
 - Subdivides vertices into **levels** L_i
- BFS on a graph with n vertices and m edges takes $O(n + m)$ time
- BFS can be further extended to solve other graph problems
 - Determines whether G is **connected**
 - Computes the **connected components** of G
 - Computes a **spanning forest** of G
 - Find and report a **shortest path** (minimum # of edges) between two vertices
 - Find a simple **cycle**, if there is one

BFS Algorithm

- The algorithm uses a mechanism for setting and getting “labels” of vertices and edges

Algorithm *BFS*(*G*)

Input graph *G*

Output labeling of the edges
and partition of the
vertices of *G*

```
for all u ∈ G.vertices()  
    setLabel(u, UNEXPLORED)  
for all e ∈ G.edges()  
    setLabel(e, UNEXPLORED)  
for all v ∈ G.vertices()  
    if getLabel(v) = UNEXPLORED  
        BFS(G, v)
```




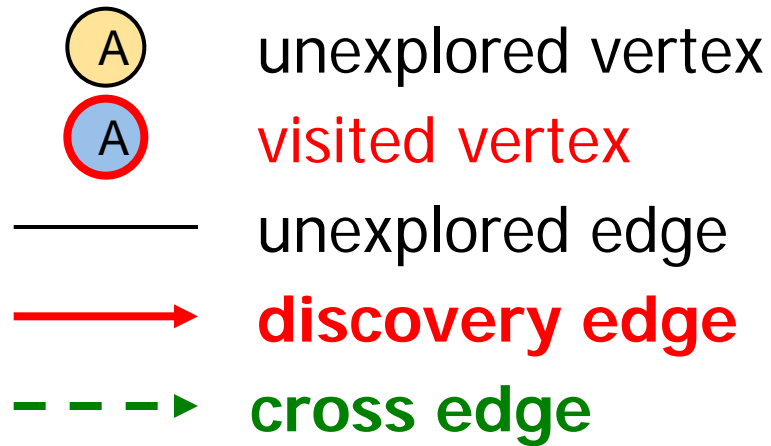
Algorithm *BFS*(*G*, *s*)

```
L0 ← new empty sequence  
L0.addLast(s)  
setLabel(s, VISITED)  
i ← 0  
while ¬Li.isEmpty()  
    Li+1 ← new empty sequence  
    for all v ∈ Li.elements()  
        for all e ∈ G.incidentEdges(v)  
            if getLabel(e) = UNEXPLORED  
                w ← opposite(v, e)  
                if getLabel(w) = UNEXPLORED  
                    setLabel(e, DISCOVERY)  
                    setLabel(w, VISITED)  
                    Li+1.addLast(w)  
                else  
                    setLabel(e, CROSS)  
    i ← i + 1
```

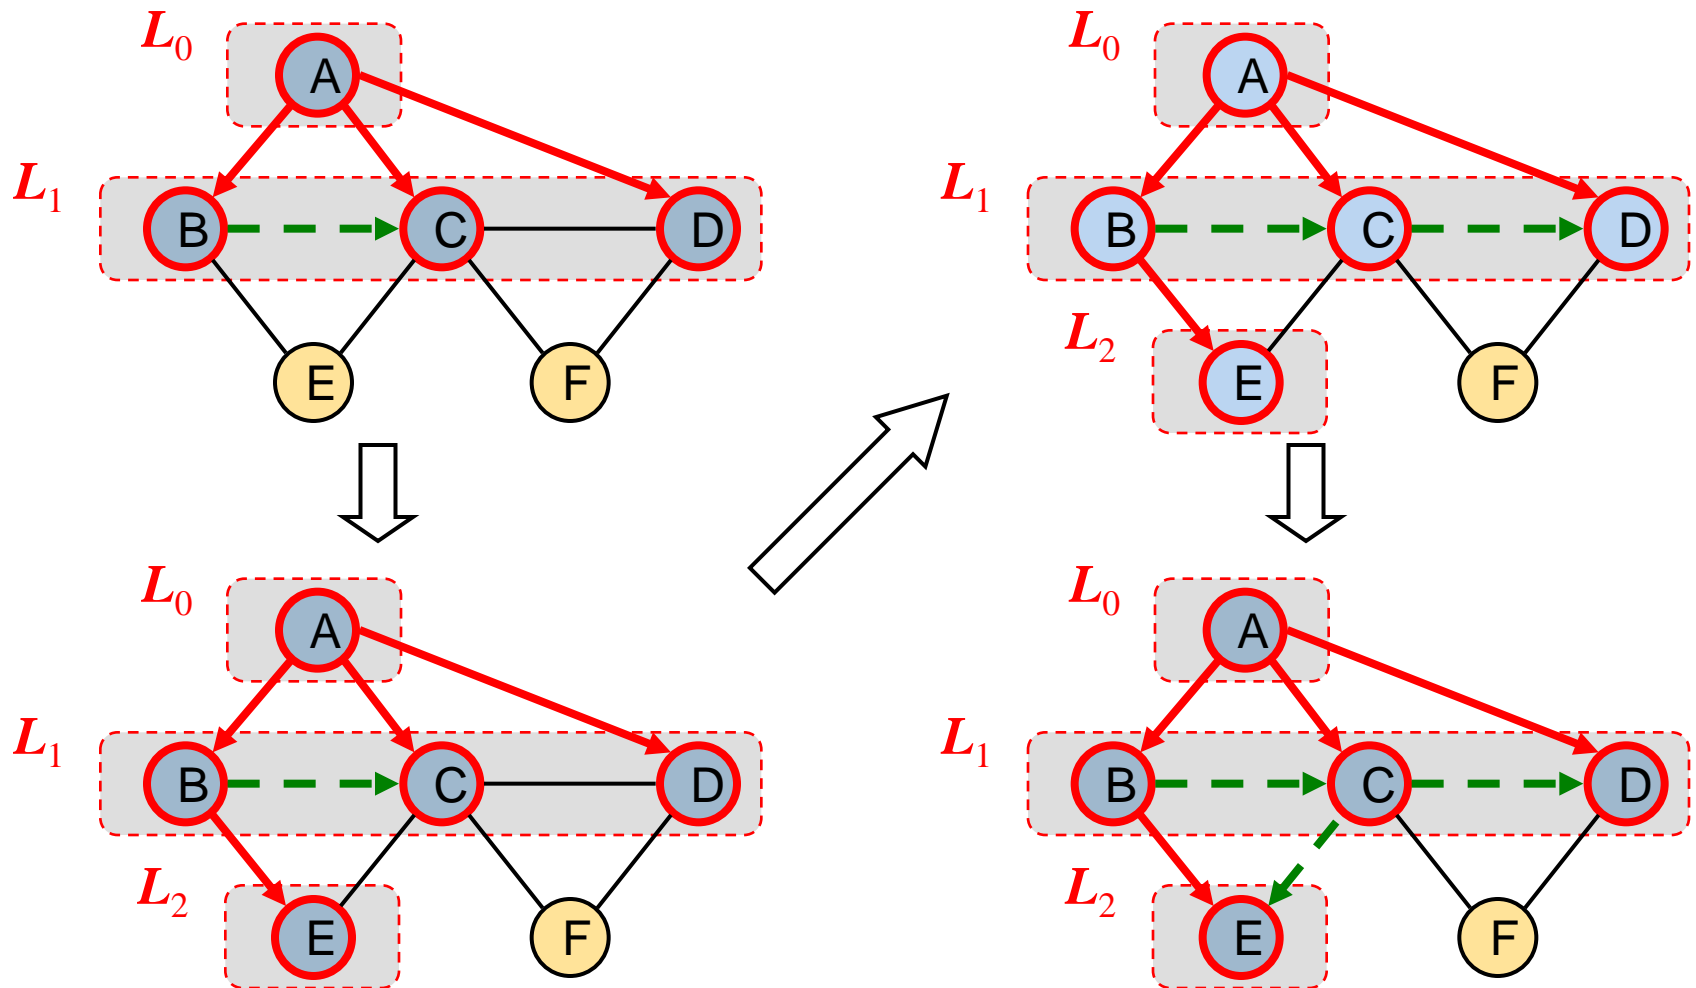
Java Implementation (fragment 14.8)

```
1  /** Performs breadth-first search of Graph g starting at Vertex u. */
2  public static <V,E> void BFS(Graph<V,E> g, Vertex<V> s,
3      Set<Vertex<V>> known, Map<Vertex<V>,Edge<E>> forest) {
4      PositionalList<Vertex<V>> level = new LinkedPositionalList<>();
5      known.add(s);
6      level.addLast(s); // first level includes only s
7      while (!level.isEmpty()) {
8          PositionalList<Vertex<V>> nextLevel = new LinkedPositionalList<>();
9          for (Vertex<V> u : level)
10             for (Edge<E> e : g.outgoingEdges(u)) {
11                 Vertex<V> v = g.opposite(u, e);
12                 if (!known.contains(v)) {
13                     known.add(v);
14                     forest.put(v, e); // e is the tree edge that discovered v
15                     nextLevel.addLast(v); // v will be further considered in next pass
16                 }
17             }
18             level = nextLevel; // relabel 'next' level to become the current
19         }
20     }
```


Example



Example (cont.)



Example (cont.)




Properties

Notation

G_s : connected component of s

Property 1

$BFS(G, s)$ visits all the vertices and edges of G_s




Property 2

The **discovery edges** labeled by $BFS(G, s)$ form a **spanning tree** T_s of G_s

Property 3

For each vertex v in L_i

- The path of T_s from s to v has i edges (**shortest path**)
- Every path from s to v in G_s has at least i edges



Analysis


- Setting/getting a vertex/edge label takes $O(1)$ time
- Each vertex is labeled twice
 - once as UNEXPLORED
 - once as VISITED
- Each edge is labeled twice
 - once as UNEXPLORED
 - once as DISCOVERY or CROSS
- Each vertex is inserted once into a sequence L_i
- Method incidentEdges is called once for each vertex
- BFS runs in $O(n + m)$ time provided the graph is represented by the adjacency list structure
 - Recall that $\sum_v \deg(v) = 2m$

Applications


- we can specialise the BFS traversal of a graph G to solve the following problems in $O(n + m)$ time
- using the template method pattern (set, map): similar to DFS
 - Compute the **connected components** of G
 - Compute a **spanning forest** of G
 - Find a simple **cycle** in G , or report that G is a forest
 - Find a **shortest path** (minimum number of edges) between two vertices, or report that no such path exists

DFS vs. BFS

Applications	DFS	BFS
Spanning forest, connected components, paths, cycles (undirected graph)	✓	✓
Shortest paths		✓
Strongly-connected components Directed cycle (directed graph)	✓	



DFS




BFS

DFS vs. BFS (cont.)

Back edge (v, w)


- w is an ancestor of v in the tree of discovery edges



DFS

Cross edge (v, w)

- w is in the same level as v or in the next level



BFS

Outline

– **Graphs (section 14.1)**

1. Definitions
2. Graph ADT

– **Data structures for graphs (section 14.2)**

1. Edge list structure
2. Adjacency list structure
3. Adjacency map structure
4. Adjacency matrix

– **Graphs traversals (section 14.3)**

1. DFS
2. BFS
3. Applications